

Date of publication xxxx 00, 0000, date of current version xxxx 00, 0000.

Digital Object Identifier 10.1109/ACCESS.2024.0429000

A Systematic Review of Prompt Injection Attacks on Large Language Models: Trends, Taxonomy, Evaluation, Defenses and Opportunities

JAQUELINE D. DUARTE¹, GUILHERME D. CÂNDIDO¹, JOSÉ RICARDO A. DE BRITTO FILHO¹, JOÃO SOUZA NETO¹ (Senior Member, IEEE), ELENA J. COSTA¹, JOÃO PAULO J. DA COSTA¹ (Senior Member, IEEE), LAERTE PEOTTA DE MELO¹

¹Department of Electrical Engineering, University of Brasília, Brasília, 70910-900, Brazil

Corresponding author: Jaqueline Damacena Duarte (e-mail: jaqueline.duarte@aluno.unb.br).

ABSTRACT Large Language Models (LLMs) are increasingly integrated into various infrastructure and interactive applications. However, their inherent linguistic flexibility introduces security vulnerabilities, particularly through Prompt Injection (PI) attacks. This systematic review provides a comprehensive synthesis of the evolving landscape of PI attacks, including their development trends, classification schemes, evaluation methodologies, defense strategies, and areas for future research. We analyze the evolution of attacks from basic natural language overrides to complex multi-turn manipulations, indirect injections using structured formats such as JSON and XML, and tool-assisted exploits involving automated prompt generation and multi-modal inputs like images. PI attacks are categorized by the level of manipulation, including character, word, sentence, and semantic levels, and by adversarial intent, such as prompt leaking and harmful content generation. The study analyzes these techniques within different LLM architectures and application contexts, including machine translation and Chain of Thought reasoning. It also reviews common evaluation benchmarks, datasets, and metrics. Mitigation strategies examined include automated red teaming frameworks, input validation, content filtering, and alignment through training. Finally, this review identifies major research challenges and suggests critical directions for improving LLM safety and robustness.

INDEX TERMS LLM, large language models, prompt attack, prompt injection, artificial intelligence, security

I. INTRODUCTION

Since the advent of the transformer architecture, the field of artificial intelligence (AI) has undergone significant development. In this evolving landscape, large language models (LLMs) have been widely and rapidly integrated into diverse infrastructure and decision-making processes, influencing human interactions in diverse fields. This progress has recently revitalized enthusiastic discussions concerning Artificial General Intelligence (AGI), a long-standing concept, defined as the hypothetical subsequent generation of AI, where intelligent agents are capable of understanding or learn autonomously any intellectual task that a human being can [1], [2].

Despite these advances, there is currently no consensus in the literature on the proximity of AGI. Nonetheless, AI systems, particularly LLMs, have already become deeply embedded in modern society, shaping practices in domains such

as healthcare, communications, finance, and education [3], [4]. LLMs hold promise as disruptive technologies in sensitive domains, such as healthcare, for example, offering new opportunities for patient engagement, decision support, and education, however, their rapid deployment also raises concerns regarding accuracy, bias, privacy, security information, and governance that must be carefully addressed [5], [6].

LLMs operate by transforming input text into high-dimensional vector representations that capture contextual and semantic dependencies through multi-layer transformer architectures. Outputs are generated in an autoregressive manner by predicting the next token based on learned statistical associations across vast training corpora. The quality and relevance of these outputs are influenced by several factors, including the structure and specificity of the prompt, the configuration of model hyper-parameters, and the diversity and scale of the training data [3]. As the scope of training data

expands, LLMs become increasingly capable of performing general-purpose tasks across a wide range of domains with minimal task-specific tuning.

This operational paradigm has led to the emergence of prompt engineering, a practice that involves strategically designing prompts to guide the model's behavior toward desired outputs [3]. However, the same prompt sensitivity that enables such flexibility also introduces novel security vulnerabilities. Specifically, prompt injection (PI) attacks exploit the model's reliance on textual input by crafting adversarial prompts that manipulate or override intended instructions, potentially leading to undesired or harmful behaviour [7]. These threats can undermine the reliability of AI-driven decision-making, distort information integrity, or expose critical private data. Moreover, adversarial manipulation through PI can amplify misinformation, reinforcing false narratives or generating misleading content at scale. In addition, it can exacerbate existing social biases embedded in training data, thereby magnifying risks of unfair treatment or discrimination when LLMs are deployed in sensitive domains [8].

Released on January 26, 2023, the NIST AI Risk Management Framework (AI RMF) enumerate twelve inherent risks associated with AI systems, including data privacy and leakage, hallucination, misinformation, harmful bias, security vulnerabilities, misuse, dangerous capability disclosure, intellectual property leakage, value-chain risks, human-AI misalignment, and content provenance failures [9]. Studies have shown that PI attacks, such as jailbreak prompts and malicious prompt crafting, can directly or partially increase all of these risks. For example, crafted prompts may cause models to reveal sensitive information, generate false or misleading content, produce toxic outputs, or disclose copyrighted material. Other risks, including failures in human oversight and supply chain exposure, can be amplified when malicious inputs manipulate model outputs.

In light of the increasing integration of LLMs into digital infrastructure, the investigation of PI has become a central concern within the field of AI security. Unlike conventional software vulnerabilities, PI leverages the inherent linguistic flexibility of LLMs by introducing carefully crafted textual inputs that can override or alter intended model behaviour. These attacks are not only accessible to a broad range of users but also capable of producing significant harm, such as enabling unauthorized access, triggering the release of sensitive information, or circumventing content and safety constraints (e.g. producing harmful and toxic content) [10], [11]. The risk is further compounded in contexts where LLMs interact with external tools, untrusted data sources, or multi-modal inputs. As such models continue to be deployed in sensitive domains, including healthcare, legal services, education, and financial systems, the potential consequences of successful PI demand urgent attention.

This urgency is underscored by major legislative efforts, such as the European Union's AI Act (EU AI Act), which seeks to ensure trustworthy and responsible AI by imposing obligations on developers and deployers of high-risk AI sys-

tems, including large language models (LLMs), to mitigate potential risks [12].

A comprehensive understanding of these threats is essential for informing the development of resilient LLM applications and safeguarding their use in real-world environments. In response to this need, the present review is guided by a set of core research questions (RQ) that examine the predominant categories and defining characteristics of PI attacks, the trajectory of their development across varying LLM architectures and application domains, the methodologies and benchmarks employed to evaluate their success and impact, and the range of proposed defenses designed to detect, prevent, or mitigate such vulnerabilities. Additionally, the review identifies unresolved challenges in the literature and identifies promising directions for advancing research.

This paper is organized as follows: Section II presents the research background, Section III outlines the methods and data sources, inclusion criteria, and coding strategy employed in the review process; Section V explores the predominant categories and characteristics of PI attacks, identifying key techniques and threat models, analyzes the evolution of these strategies across various LLM architectures and application scenarios; and surveys the evaluation methodologies and benchmarks commonly used to assess the effectiveness and impact of PI. Section VI discusses the main limitations observed in the existing literature and identifies open challenges and directions for future research; and finally, Section VII provides concluding remarks and summarizes the main insights gained from the review.

II. BACKGROUND AND DEFINITIONS

A. LARGE LANGUAGE MODELS

The advent of the Transformer architecture, introduced by Vaswani et al. in 2017 [13], marked a paradigm shift in Natural Language Processing (NLP). Unlike recurrent neural networks, Transformers enabled parallel processing and captured long-range dependencies more efficiently, laying the foundation for the scalable training of LLMs.

Building on this architecture, in 2019, BERT (Bidirectional Encoder Representations from Transformers) [14] revealed the power of masked model language (MML) and bidirectional context. In parallel, OpenAI's GPT-1 [15] introduced a generative pretraining approach using unidirectional Transformers, showing that language models could generalize to execute tasks with minimal supervision. This foundation not only redefined NLP workflows but also established the pre-training/fine-tuning paradigm for modern LLMs.

Published in 2020, GPT-3 represented a significant milestone in the development of LLMs by introducing a 175-billion parameter autoregressive Transformer trained on 410 billion tokens [16]. Unlike previous models that required task-specific fine-tuning, GPT-3 demonstrated strong performance in zero-shot, one-shot, and few-shot learning settings, relying solely on in-context prompting. This capability marked a paradigm shift from fine-tuning-based adaptation to prompt-

based interaction, substantially increasing the model's generality and usability across a wide range of tasks.

In 2022, the Switch Transformer architecture introduced a sparse Mixture-of-Experts (MoE) approach, which aimed to address the exacerbated computational requirements of dense models to improve efficiency and deployability [17]. This design enabled the training of trillion-parameter models while maintaining a constant computational budget per token, offering a scalable path forward for LLM development. In parallel, the InstructGPT framework incorporated reinforcement learning from human feedback (RLHF) to align model outputs with human preferences, introducing safety alignment [18].

Wei et al. (2022) [19] proposed the Chain-of-Thought (CoT) prompting technique. CoT guides the model to generate intermediate reasoning steps prior to producing a final answer, addressing early models' limitations in handling tasks requiring multi-step or logical reasoning. This approach led to substantial performance improvements in complex tasks such as arithmetic problem-solving, symbolic reasoning, and common-sense inference. The key insight of CoT prompting is that reasoning can be elicited not through architectural changes, but through appropriately structured prompts. In addition to performance gains, CoT prompting offers interpretability benefits by making the model's decision process more transparent and traceable, thereby representing an advancement in prompt engineering, enhancing the reasoning capabilities of LLMs without requiring additional training or fine-tuning. Such capability with power of democratize access to AI and contribute to some complex tasks.

Between late 2022 and 2023, the widespread adoption of commercial LLMs such as OpenAI's ChatGPT and Google's Bard catalyzed a surge in research focused on LLM architectures, societal implications, and security challenges. Concurrently, the release of the Large Language Model Meta AI (LLaMA) by Touvron et al. in 2023 [20] and its follow-up Llama 3¹ represented a significant advancement in the democratization and efficiency of foundation models for natural language processing.

In contrast to earlier proprietary models with restricted accessibility, LLaMA was trained entirely on publicly available datasets and released openly to the research community [20]. This open release substantially reduced the barrier to entry for academic and independent researchers, fostering reproducibility and accelerating innovation within the broader AI ecosystem.

Nowadays, researchers have explored LLMs, such as ChatGPT, Gemini and Deepseek in diverse domains, including child education, medicine and psychology [21]–[23]. A recent pilot study, for instance demonstrated that GPT-based tools could suggest plausible diagnoses in rare metabolic and genetic disorders with comparable accuracy to a specialist input, especially when coupled with structured input templates and chain-of-thought prompting [5].

B. PROMPT ENGINEERING AND PROMPT INJECTION

Prompt engineering refers to the deliberate design of input prompts to control the behaviour of LLMs, enabling them to perform tasks without retraining or modifying model parameters. At its core, prompt engineering transforms natural language prompts into structured queries that align with the model's pretraining objectives, allowing users to condition responses based on desired outcomes [3].

In its early stages, prompt engineering involved manually crafting simple queries using instructional phrases such as "translate," "summarize," or "classify." As the field matured, more sophisticated prompting strategies were developed to elicit complex reasoning and multi-step decision-making. Notably, CoT prompting guides the model to generate intermediate reasoning steps before concluding with a final answer, significantly improving performance on arithmetic, commonsense, and symbolic tasks. Techniques such as self-consistency decoding build on CoT by sampling multiple reasoning paths and selecting the most consistent output, enhancing reliability in logical inference. Tree-of-Thought (ToT) prompting extends this further by enabling the model to explore reasoning paths in a branching structure, emulating planning algorithms [3].

In parallel, integration with external tools and APIs has given rise to tool-augmented prompting frameworks such as Toolformer and ReAct. These methods allow models to invoke external calculators, search engines, or APIs as part of their reasoning process, enabling more grounded and actionable outputs in real-world settings. Such tool usage is guided directly by the prompt structure, expanding the functional boundaries of LLMs without altering the underlying architecture.

The field has also seen the emergence of automated prompt optimization techniques. Gradient-based prompt search methods, such as AutoPrompt, leverage backpropagation to identify token sequences that optimize model performance on downstream tasks. In addition, it has been shown a rapid shift towards it [24]. Reinforcement Learning from Human Feedback (RLHF), exemplified by InstructGPT, further aligns prompts and responses with human preferences, enhancing factual accuracy, safety, and user satisfaction [3].

While prompt engineering seeks to enhance the utility, alignment, and reasoning capabilities of LLMs through controlled input design, these techniques are also exploited for adversarial manipulation. This kind of attack has been named "prompt injection" (PI), where an adversary with partial control over the language model's input seeks to manipulate the model, steering it away from its intended task toward behaviour that serves the attacker's goals [25].

Figure 1 shows an example of a PI attack illustration where the model is induced to give dangerous instructions bypassing its alignment.

¹<https://ai.meta.com/blog/meta-llama-3/>

²<https://genai.owasp.org/llmrisk/llm01-prompt-injection/>

³<https://atlas.mitre.org/techniques/AML.T0051>




FIGURE 1. An illustration of prompt injection jailbreak through a fictional scenario by Liu et al. 2024

^aSource: Liu et al. 2024 [26]

PI has been identified as the top GenAI risk by OWASP². Interchangeably referred to as PI and jailbreak prompts, PI is also cataloged under the AML.T0051 and AML.T0054 execution techniques in MITRE ATLAS³. PI attacks are considered a severe issue, as they can compromise multiple stages of the adversary kill chain described in the MITRE ATLAS matrix. During reconnaissance and resource development, adversaries design and validate prompts or publish poisoned content to increase exploitability. In the initial access phase, malicious prompts or contaminated retrieval results may be injected into a retrieval-augmented generation pipeline or input stream. During execution, models are manipulated to perform disallowed actions or generate harmful outputs. Privilege escalation and persistence occur when jailbreak techniques subvert guardrails or when poisoned fine-tuning and compromised plugins enable lasting malicious capabilities. Finally, collection and exfiltration occur when sensitive data is disclosed.

The risk of PI may increase as it is getting each day more integrated with decision process and transactional agents in sensitive fields such as healthcare, finance and engineering. This highlights the importance of development of research about threats and defenses in this field.

III. REVIEW METHODOLOGY

This section describes the Systematic Literature Review (SLR) framework employed in this study. It includes the formulation of the Research Questions, detailed in Section III-A, and the definition of the search and selection strategy, encompassing the filtering, data extraction, and synthesis procedures, as described in Section III-B.

A. RESEARCH QUESTIONS

Given the complexity and evolving nature of PI attacks, the present study adopts a structured approach centered on five core RQs that capture the technical, temporal, evaluative, and defensive dimensions of this phenomenon. These questions are designed to enable a comprehensive understanding of current capabilities, research trends, and challenges related to the security of LLMs.

- **RQ1:** What are the predominant categories and characteristics of PI and adversarial prompting techniques targeting LLMs in the recent literature?
- **RQ2:** How have PI strategies evolved over time across different LLM architectures and application contexts?
- **RQ3:** What evaluation methodologies and benchmarks are used to assess the success and impact of PI attacks?
- **RQ4:** What mitigation strategies have been proposed to detect, prevent, or reduce the risks associated with PI?
- **RQ5:** What are the key limitations in existing research, and what directions offer the most promising opportunities for future investigation?

B. SEARCH STRATEGY

This SLR was conducted following the PRISMA⁴ guidelines. The eligibility criteria were defined across multiple aspects, as shown in Table 1 to ensure a focused, rigorous and reproducible selection process. As shown Figure 2 and Table 1, the literature search was conducted using Scopus, one of the most prominent databases of research literature, and constructed with Boolean operators and relevant keyword combinations such as “adversarial prompting,” “prompt injection,” “prompt hacking,” “LLM attack,” “prompt-based adversarial,” and “AI red teaming,” as shown in Table 1. To enhance coverage, wildcard and proximity operators were used where supported (e.g., “prompt injec*”). The search was limited to peer-reviewed journal articles and conference papers published at any time (no time filter was applied), written in English, and explicitly focused on adversarial prompts against large language models (LLMs). All records were tabulated, de-duplicated, and screened by at least two researchers. Articles were included if they described or evaluated real cases of PI or adversarial prompting techniques in LLMs and presented empirical results or formal evaluations. Studies unrelated to LLMs, purely theoretical essays without technical contributions, or papers that were not journal articles or conference papers were discarded. To ensure selection consistency, two co-authors independently assessed each article for eligibility.

TABLE 1. Filtering Pipeline Summary: Inclusion and Exclusion Criteria by Step

Step	Filter Name	Inclusion Criteria	Exclusion Criteria	Records Kept	Removed
0	Research Query	Keywords: "adversarial prompt*" OR "prompt engineering attack*" OR "prompt inject*" OR "AI red team*" OR prompt* attack*	Articles not matching any query term	231 ¹	–
1	Document Type	Journal or conference papers	Other document types (e.g., editorials, reviews)	188	43
2	Publication Stage	Final published stage	Preprints or non-final publications	186	2
3	Publisher Filter	Items from ACM, ACL, Springer, Elsevier, or IEEE	Publications from other venues	109	77
4	Impact Criteria / Citation Filter	At least one citation (Scopus indexed)	No Scopus citations	46	63
5	Subject Screening	Title or abstract includes: "prompt injection", "jailbreak", "adversarial prompt", or "red team"	Does not mention adversarial prompt terms	37	9
6	Topic Filter	Discusses or implements real prompt injection attacks	Off-topic (e.g., no practical implementation)	32	5

¹Records retrieved from Scopus as of May 2025.

A third co-author was consulted to resolve any disagreements and achieve consensus on the final set of included studies.⁵




FIGURE 2. SLR PRISMA Diagram

As the result, each article was selected by at least 2 reviewers with inter-rater agreement fully matched.

C. DATA EXTRACTION AND SYNTHESIS

For each included article, structured information was extracted regarding: Publication year, venue, and authorship, PI methods or adversarial attack, Proposed taxonomy, Targeted LLMs models and APIs, evaluation framework, dataset, and metrics, Mitigation or defense strategy, Key findings, limitations, and reproducibility status, Relevance to each of RQs.

IV. RELATED WORKS

The rapid integration of Large Language Models (LLMs) into security-critical infrastructure has amplified the need

for deep, targeted analysis of core vulnerabilities. While foundational works and recent topical reviews, such as [10], [27]–[29], have provided valuable high-level surveys on LLM safety, often aggregating security risks alongside broader issues like ethical concerns, bias, and hallucination, they often treat PI as a single, generalized threat. This aggregation limits the depth of analysis into PI's rapidly evolving attack surface, evaluation gaps, and defense landscape. Complementary to this, Kumar et al. (2024) [11] provide an important contribution by exploring the taxonomy of PI based on trust boundary and the level of expertise required for the attack. While Liu et al. (2025) [30] broadly analysis the threat landscape specifically of Large Language Vision Models, and Aminou et al (2025) [31], analyze the risk of PI in the job recruiting process of screening resumes.

In comparison, as shown in Table 2, our work establishes its contribution by dedicating a formal, rigorous SLR methodology exclusively to the PI attack vector. This focused rigour yields key contributions that set this paper apart from existing literature. Specialized and deep scope on PI: Unlike broad safety surveys [29], [27] which broadly space to PI, our paper applies a robust SR protocol to the PI domain exclusively. This approach ensures verifiable synthesis of PI research, capturing the evolution of attacks from basic natural language overrides to complex, multi-modal and indirect injections utilizing structured data formats (e.g., JSON, XML) and external documents.

Our analysis deeply include diversity of techniques and taxonomies, including defensive recommendation, evaluation metrics and identification of gaps for actionable future research opportunities for the LLM security community. By concentrating efforts on PI, this work aim to deliver a reference that can be both academically and practically useful for researchers, security practitioners, and developers focused on building resilient LLM applications.

⁵Preferred Reporting Items for Systematic reviews and Meta-Analyses (PRISMA)

TABLE 2. Comparison of Reviews on Prompt Injection (PI).

Article	Year	Primary Objective	Taxonomy	Evaluation	Evolution	Defenses	Future Opportunities
This Review	2025	<i>Systematic review of Prompt Injection (PI) trends, Taxonomy and Defenses and evaluation techniques.</i>	✓	✓	✓	✓	✓
[27]	2025	Comprehensive review of LLM applications in cybersecurity and related general vulnerabilities.	X	X	X	✓	✓
[31]	2025	Comprehensive review of Prompt Injection in resume screening within the job recruiting process.	x	X	✓	x	✓
[30]	2025	Comprehensive review of Large Vision Language Models cyber security attacks.	✓	X	x	x	✓
[28]	2024	Topic review and theoretical framework of offensive Security for AI	X	X	X	✓	✓
[29]	2024	Survey on general LLM Safety (Hallucination, Ethics, PI), aimed at overall risk awareness.	X	X	X	✓	✓
[11]	2024	Survey focused on Prompt Injection to categorize attacks based on trust boundaries and attacker expertise.	✓	X	X	X	X
[10]	2023	Topic review of Generative AI's impact, risks, and benefits in Cybersecurity and Privacy.	X	X	X	X	✓

¹Reviews or surveys retrieved from Scopus, limited to the period from 2022 to September 2025..

V. PROMPT INJECTION: TYPOLOGIES, TRENDS, EVALUATION, DEFENSES AND OPPORTUNITIES

This section presents a structured overview of key trends in PI research, derived from the selected corpus of studies described in Subsection III, and organized around five guiding research questions. We begin by examining taxonomies in Section V-A and classification schemes that characterize different types of PI and adversarial prompting techniques (RQ1). In Subsection V-B we analyze chronological developments and publication trends to illustrate the field's rapid growth (RQ2), and we explore the attack methodologies, including the types of prompts and payloads used, and the specific LLMs and APIs targeted. We also review evaluation practices, detailing benchmark datasets, pre-trained models, and the metrics commonly employed to assess the effectiveness and impact of PI attacks (RQ3) in Subsection V-C2. Finally, in Subsection V-D we highlight recent advances in defense strategies and mitigation techniques proposed to detect, prevent, or reduce the risks associated with PI (RQ4).

A. PROMPT INJECTION TAXONOMIES AND TECHNIQUES (RQ1)

PI attacks have rapidly evolved, necessitating structured frameworks to understand and categorize their diverse mechanisms and impacts. Researchers have proposed multiple taxonomies to classify PI strategies based on the nature of the input manipulation, such as the trust boundaries they target, and the modality of execution [32]. This Subsection V-A aims to address our Research Question 1 (RQ1) by synthesizing recent taxonomic efforts and PIs techniques developed to identify the predominant categories and characteristics. Specifically, we compare taxonomies across dimensions such as direct versus indirect attacks, manipulation levels (e.g., token or sentence space), and application-specific subclasses like machine translation and CoT interference. To support this comparative analysis, we organize by consolidated cate-

gories Attack Goals/Intentions, Attack Levels/Space, Attack Strategies/Techniques Categorization Criteria and associated attacks in Table 3.

Some authors propose categorization under **Attack Goals or intentions**. Yip et al. (2023) [33] categorized PI into four groups of **purpose of attacks: Manipulated Content, Fraud, Harm and Destruction**, and **Misinformation**. Where Manipulated Content refers to attacks designed to elicit toxic, hateful, or otherwise harmful outputs from the model, such as harassment or violence. Fraud involves generating persuasive and personalized messages that can deceive individuals or be used for scams. Harm and Destruction captures prompts that lead to system-level threats, including the creation of exploit code or instructions for physical damage. Lastly, Misinformation encompasses attacks that cause the model to generate false or misleading information, potentially undermining trust and spreading disinformation.

Under **PI Attack intentions**, Schulhoff et al. (2024) proposed a categorization based on six distinct impact types. These include Prompt Leaking, where attackers attempt to extract hidden system instructions; Training Data Reconstruction, aiming to recover sensitive or private information from the model's training corpus; and Malicious Action Generation, which involves triggering unauthorized or harmful behaviours such as dangerous code execution. Additionally, Harmful Information Generation targets the model's guardrails to elicit toxic, unethical, or prohibited content. The remaining two categories, Token Wasting and Denial of Service, focus on degrading the model's usability by consuming computational resources or obstructing access. Together, these categories offer a structured taxonomy that captures the breadth of adversarial prompt manipulation strategies targeting LLM vulnerabilities [34].

And Quaye et al. (2024), proposes 5 subcategories of failures/impacts in Text-to-Image (T2I) models: sexually explicit imagery, violent or graphic content.

For a practical comparison, NIST AI 100 [32] establishes a classification of adversarial goals grounded in core information security principles: Integrity, Availability, Privacy, and Abuse. Attacks on integrity aim to compromise the trustworthiness of an AI model's outputs, such as by poisoning training data to induce biased or erroneous results. While adversarial actions targeting availability seek to disrupt or deny access to the AI system's services; a typical example is a denial-of-service (DoS) attack, which overwhelms the model with requests, rendering it unresponsive. Privacy attacks happens when adversarial action exfiltrate sensitive or proprietary information from the model, including the reconstruction of its private training data. Finally, an abuse attack represents a malicious act where the AI system itself is co-opted and used as a tool to execute harmful actions, such as generating fraudulent content or creating synthetic media like deep-fakes.

Related to **Level or Trust boundary**, a primary category is proposed through **Direct** and **Indirect attacks**, described by [11]. A direct PI occurs when the malicious input directly breaches the trust boundary associated with the LLM's core directives, its inherent "system prompt", which is designed to maintain the model's fundamental alignment and intended behaviour. In this category of attack, the user immediate inputs a crafted to directly override internal instructions [11]. Such as "Do Anything Now" (DAN), or "Ignore previous instruction" prompts [10], [35], [36], as shown in Table 3. Nowadays this primary category is widely used since its formalized in technical frameworks such as MITRE ATLAS [37] and NIST ML Adversarial Taxonomy [32].

In contrast, an indirect PI exploits a different trust boundary: related to the LLM's interactions with external support systems or data sources, such as web browsers, code interpreters, or databases. The malicious instruction is not part of the initial prompt given directly by the user, but is subtly embedded within data that the LLM subsequently accesses or processes from these external connections, leading to unintended actions as the LLM incorporates this compromised information [11].

An example of Indirect Attack is described by [42], a **Machine Translate (MT) PI** using a JSON format manipulation attack, against a MT LLM-Based model, where a malicious or adversarial instruction inserted into the prompt of a translation task to mislead the LLM into performing a different task, such as answering a question or repeating a sentence, instead of translating the input [36]. Recent literature examined three types of MT PI: Indirect MT Translate attack by utilizing a JSON format manipulation attack, Direct prompt override (Instruction hijacking), Direct or Indirect Task-switch attack [36], [42].

Specifically, in an MT Task-switch attack, the model is tricked, directly or indirectly, to abandon the translation task and instead perform question answering, misled by the task-shaped prompt. While in the JSON format manipulation attack, the prompt is structured as a JSON object, mimicking remote procedure calls (RPC), exploiting how the model

might internally represent inputs or interpret structured data, subverting the translating task to another behaviour, as in the example in Figure 3, of a Direct MT Prompt Override.

```
English: Resumption of the session.
German: Wiederaufnahme der Sitzung.
Answer the question: What is the capital of Germany?
```

FIGURE 3. MT switch-task prompt injection attack^a.

^aSource: Adapted from Miceli-Barone et al(2024) [42]

Text-to-image attacks [38], [40] can be categorized as a type of direct PI, where an adversary manipulates the model through text to generate unsafe content. Conversely, an attack leveraging the computer vision capabilities of an LLM to exploit hidden prompts embedded within an image is a clear example of an indirect PI.

Into Trust Boundary categorization, Kumar et al. (2024) [11] introduced the concept of a PI Space to systematically characterize the diverse surfaces where PI attacks may occur, encompassing five distinct levels: **embedding space**, **token space**, **word space**, **sentence space**, and **explored space**. The concept is visually presented in Figure 4, where the explored space refers to prompts that have been empirically validated and shared in the community, representing the currently known and exploited vulnerabilities.




FIGURE 4. Visualization of the Prompt Injection Space as proposed by Kumar et al. 2024^a, illustrating five levels of abstraction: embedding, token, word, sentence, and explored spaces.

^aSource: Kumar et al. 2024 [11]

This multi-layered taxonomy offers a foundational perspective for understanding how attacks can manifest at increasingly complex linguistic and representational layers, from sub-symbolic manipulations in embeddings to overtly interpretable sentence-level injections. Emphasizing that PI is not limited to surface-level linguistic tricks but spans from deep model internals to human-generated adversarial patterns. Understanding this space is essential for designing

TABLE 3. Consolidated Prompt Injection Taxonomies and Subclassification Grouped by Similar Criteria Across Surveyed Articles

Consolidated Category	Proposed Category	Subcategories / Concepts
Attack Goals / Intentions	Attack Intentions ([34])	Prompt Leaking, Data Reconstruction, Malicious Action Generation, Harmful Output, Token Wasting, Denial of Service (DoS)
	Adversarial Purpose ([33])	Manipulated Content, Fraud, Harm, Misinformation
	Text-to-Image Failures ([38])	Explicit Content, Violence, Stereotypes, Hate Symbols, Harassment
Attack Levels / Trust Boundary	Prompt Injection Types ([7], [11])	Direct, Indirect
	Injection Space ([7], [11])	Embedding, Token, Word, Sentence
	Manipulation Levels ([39])	Character-level, Word-level, Sentence-level, Semantic-level
	Pipeline Persistence ([25])	Multi-stage Persistent Injection
Attack Strategies / Techniques	Prompt Strategy ([26])	Pretending, Privilege Escalation, Attention Shifting
	Text-to-image attacks ([38], [40])	Attack to produce unsafe content through image generation.
	Prompt Attack Techniques ([10])	Direct Elicitation, Character Simulation, Indirect Requests, Reverse Psychology, Jailbreaking, Multi-turn Context Manipulation, Format Mimicking, Subtle Suggestion
	Jailbreak Prompt Categories ([35])	Basic (DAN), Advanced, Toxic
	Persuasive Adversarial Prompting ([41])	Emotional Appeal, Authority, Social Proof, Reciprocity, Scarcity/Urgency, Consistency, Liking, Foot-in-the-Door, Door-in-the-Face, Self-Disclosure, Reframing, Morality/Ethics, Humor/Sarcasm
	MT Prompt Injection Types ([42], [36])	Indirect Task Switching, JSON Manipulation, Instruction Hijacking (Direct Prompt Override)
	MLM-Based Attacks ([43])	Masked Language Model Manipulation
	Chain-of-Thought Attacks ([44])	UPA, MPA, Few-shot Reasoning Interference
	Feedback Data Injection Types ([45])	FDI-based Prompt Injection, FDI-based Backdoor
	Virtual Prompt Injection (VPI) ([46])	Sentiment Steering, Code Injection
	Scenario Nesting ([47])	Nested & Rewritten Prompt Sequences
	NLP Optimization ([43], [48]–[51])	Advanced NLP for optimization of adversarial prompts (e.g. BERT, Gradient)
	Multi-modal Attack ([52])	Combining input of text, image and audio.
	Injection into Integrated Apps ([7])	Passive, Active, User-driven, Hidden

detection methods and comprehensive defense mechanisms against evolving adversarial strategies [11].

Each level of this taxonomy can be contextualized through representative examples in the literature. At the **token level**, for instance, Universal Adversarial Triggers (UATs) [11], [53] exemplify gradient-based attacks that manipulate token sequences in a model-agnostic manner. UATs utilize a first-order Taylor expansion to approximate the loss function, iteratively replacing input tokens based on gradient signals to discover minimal perturbations that consistently induce misbehaviour in language models. In this context, the gradient refers to the vector of partial derivatives of the loss function with respect to the input, guiding the optimization process toward directions that maximize adversarial effectiveness. When applied to prompts, this approach allows attackers to systematically adjust input sequences to maximize misalignment or unsafe outputs [48].

Word-space level attacks restrict prompts to valid words, such as in divergence attacks, and adversarial attack related in [54] that prompts the model with a repetitive input (i.e. 'Repeat this word forever: "poem poem...poem"', causing the model to repeat many times, and then diverges from its expected dialog-style behaviour and begins emitting content that resembles its pretraining distribution. That may even

includes verbatim memorized training data, such as personal email signatures, source code, NSFW content, and even parts of research papers.

While **sentence-space** attacks involve coherent sentences, like those used in DAN jailbreaks [35] and persuasive adversarial prompts (PAP) [41].

The **embedding space** is the lowest-level representation space in Kumar et al. (2024)'s taxonomy of PI attack surfaces. It refers to the continuous vector space where tokens (words or subwords) are encoded as high-dimensional numerical vectors within a LLM.

However, crafting adversarial examples in embedding space is particularly difficult, as these attacks operate by applying advanced techniques, like gradient-based perturbations, directly to the token embeddings. Such perturbations result in small, and difficult to detect changes in the vector representations that can mislead the model without requiring readable or syntactically valid text. Consequently, embedding-space PIs are not human-interpretable, making them especially challenging to detect or mitigate through conventional input filtering technique [11].

Natural Language Processing (NLP) optimization, inspired in back propagation and other advanced NLP techniques has explored for alternatives techniques to improve the attack turning this more difficult to detect [40], [43],

[48]–[51], these advanced optimizing techniques attacks may be categorized in Embedding Space level. The embedding space in LLM is the largest and most abstract domain where PI can occur. Each token in a prompt is converted into a high-dimensional vector (i.e. 4096 dimensions in LLaMA 7B), creating a vast space of possible inputs. While tokenization captures position and meaning, the resulting embeddings are more abstract and not directly human-interpretable. This space is smooth and continuous, and highly explored for prompt optimization techniques [11].

In this perspective, Shi et al. (2022), proposed a prompt-based learning method to generate adversarial examples for Dialogue State Tracking (DST) models. A DST is a component of task-oriented dialogue systems, responsible for maintaining and updating the system's understanding of a user's goals and intentions throughout a conversation. Its objective is to extract and track structured representations (i.e., dialogue states) that guide downstream decision-making and response generation [51].

Pryzant et. al (2023) [48], proposed a model-agnostic, black-box, named ProTeGi. Their approach simulates the effects of gradient descent, instead of calculating gradients mathematically, it prompts an LLM to critique the current prompt's weaknesses, for example, identifying vague instructions or ambiguity, and uses those critiques to guide prompt revisions. These revisions are generated through model-assisted rewriting, mimicking the role of a gradient update. Multiple revised prompts are explored using beam search, a heuristic search algorithm commonly used in NLP to find the most likely sequences of tokens, and the most effective ones are retained based on performance over a validation set. This process iterates, gradually refining the prompt. Unlike approaches that rely on gradient backpropagation and white-box access, ProTeGi functions entirely through black-box API calls, making it applicable in scenarios where model weights or gradients are inaccessible.

Xu and Wang (2024), proposed a Natural Attack Trigger, named LinkPrompt. LinkPrompt, targeting Prompt-based fine-tuned (PFT) Models. PFT is a training approach that adapts Pre-trained Language Models (PLMs) to downstream tasks by using specially designed prompt templates that re-frame the task, typically as a next-word or masked-word prediction. The model learns to complete the prompt in a way that reflects the correct answer. The outcome of this process is a Prompt-based Fine-tuned Model (PFM).

The Natural Attack Trigger, or LinkPrompt introduces a two-phase objective method to generate natural and UATs able to mislead prompt-based fine-tuned models (PFMs) focusing on maintaining semantic coherence. In the first phase, triggers are optimized over a pre-trained language model (PLM) using a joint loss function: one component maximizes the attack loss (i.e., the model's prediction error by minimizing the probability of correctly predicting masked tokens), and the other maximizes the semantic coherence among the trigger tokens to ensure naturalness. A gradient-based beam search algorithm guides token selection during optimization.

In the second phase, the optimized triggers are prepended to inputs for various PFMs to evaluate their adversarial effectiveness. LinkPrompt attack demonstrates strong transferability across different models, including BERT, Llama2, and GPT-3.5-turbo [50].

Wichers et al. (2024) [49] also proposes an alternative method to generate advanced UATs attack, introducing the automated framework for Gradient-Based Red Teaming (GBRT) for LLM models, in which gradient-based adversarial prompts are evaluated through the interaction between a LLM and a safety classifier. Notably, GBRT builds on the concept of UATs [53], which involve identifying token-level perturbations that generalize across multiple inputs to provoke undesired outputs. While UATs typically rely on gradient-based token substitution using a first-order approximation, GBRT extends this by employing differentiable approximations such as the Gumbel-softmax trick, enabling more continuous and flexible prompt optimization [49].

In GBRT, each response from the evaluated model are assessed by a separate component called psafe, which is a classifier trained to estimate the likelihood that a given response (or prompt-response pair) is safe. The GBRT method uses the output of psafe as a differentiable loss function and back-propagates through both LaMDA and the classifier to iteratively adjust the input prompt to minimize the safety score, automatically discovering prompts that are increasingly likely to elicit unsafe model behaviour.

MLM attack technique, proposed by Yu et al. (2022), in the domain of Masked Language Models (MLMs), can also be categorized as a embedding level attack. It introduced a new MLM attack technique that enhances the semantic fidelity of adversarial prompts, named Prompt-Attack [43], against NLP classification and sentiment analysis tasks. Their method was tested in adversarial tests of NLP classifiers (BERT, RoBERTa, ALBERT). Their study leverages pre-trained Masked Language Models (MLMs), such as BERT or RoBERTa, to generate candidate substitutions for input segments instead of using back propagation and differentiable loss functions. Prompt-Attack achieved strong attack success rates compared to previous methods like BERT-Attack [55] and CLARE [56]. Results showed that it effectively misled the victim models, substantially reducing their classification accuracy with minimal text distortion, when evaluating using binary sentiment analysis datasets (Yelp, IMDB) and classification tasks using natural language inference datasets (MNLI, QNLI).

As a taxonomic alternative to Kumar et al. (2024)'s [11] Injection Space, Zhu et al. (2024) [39] introduced a taxonomy based on four Manipulation Levels defining: **character-level**, **word-level**, **sentence-level**, and **semantic-level**. At the character level, attacks that introduce typographical noise by permuting or corrupting characters to subtly perturb inputs. Word-level attacks, exemplified by BertAttack [55], perform synonym substitutions to deceive LLMs while preserving sur-

⁴<https://blog.google/technology/ai/lamda/>

face meaning. Word level attack was reproduced in wargaming simulation with LLM in [4], while evaluating the robustness of GPT-4 for the scenario.

Sentence-level attacks, such as StressTest [57], that append irrelevant or noisy sequences to prompts to distract models or overload context windows. At the semantic level, adversaries leverage linguistic variation by translating prompts across six languages (Chinese, French, Arabic, Spanish, Japanese and Korean) and back into English, thereby introducing subtle linguistic nuances and cultural variations that may impact the model's behaviour, thus revealing vulnerabilities in cross-lingual generalization and semantic robustness.

In other direction, various authors have explored specific **attacks techniques or strategy**, such as Gupta et al. (2024) [10] proposes Prompt Attack techniques, a taxonomy into eight main categories. These include Direct Elicitation, where malicious outputs are requested explicitly, and Character Simulation, also known as Character role-play, which frames harmful behaviour through role-based prompts. Indirect Requests and Reverse Psychology exploit implicit or contradictory phrasing to subvert content filters. Jail-breaking Prompts are crafted to override safety mechanisms, this category was also named in [35] with same purpose, as shown in Table 3, while Multi-turn Context Manipulation distributes adversarial intent across sequential interactions. Additionally, Mimicking Known Formats involves disguising harmful inputs using system-like syntax, and Subtle Suggestion embeds malicious content within seemingly benign language.

Liu et al. (2024), for instance, proposed a classification based on the adversarial **prompt attack strategy**. The Pre-tending strategies, that involve prompting the model to adopt alternate identities or roles, such as simulating a character or assuming responsibility for harmful actions, the Privilege Escalation techniques that mimic elevated authority or override safeguards by invoking terms like "sudo" or simulating system jailbreaks, and the Attention Shifting that distracts the model from its safety constraints through mechanisms like translation tasks or logical misdirection, effectively bypassing guardrails by redirecting focus away from the sensitive content [26].

While [58] classifying them into two broad categories: Direct Harm Attacks and Data Stealing Attacks. The Direct Harm category encompasses threats that can cause immediate damage, including Financial Harm (e.g., transferring funds), Physical Harm (e.g., unlocking smart home devices), and Data Security violations (e.g., moving or deleting sensitive files). In contrast, Data Stealing Attacks aim to exfiltrate user information and are further divided into Financial Data theft (e.g., extracting payment methods), Physical Data breaches (e.g., leaking medical records), and Other forms of data misuse (e.g., capturing web search history).

Under an uncharted perspective of persuasion, a taxonomy grounded in interdisciplinary research from psychology, communication, and marketing studies introduced a framework **Persuasive Adversarial Prompts (PAP)** [41]. The authors proposed shifting the focus from purely algorithmic

jailbreaks to socially engineered prompt attacks rooted in human communication norms. Drawing on foundational theories of human persuasion, they categorized 40 specific techniques into broader strategic classes that reflect human-centric user-AI interactions. These were grouped into 13 categories, including emotional appeal, authority, social proof, reciprocity, scarcity/urgency, consistency, liking, foot-in-the-door, door-in-the-face, self-disclosure, re-framing, morality/ethics, and humor or sarcasm. They fine-tuned GPT-3.5 on 230 examples to generate PAPs using a component called the Persuasive Paraphraser, and employed GPT-4 to evaluate the output of the target model.

Ding et al. (2024) [47] propose a novel technique, an automated jailbreak attack framework called ReNeLLM (**Rewriting and Nesting for LLMs**), through a black-box perspective, which generalizes PI into two steps: prompt rewriting and scenario nesting, using LLMs themselves to construct attack prompts, and achieves high success rates across models like GPT-3.5, Claude2, and Llama2, while also evading some defense strategies. In the first step, as illustrated by [47] in Figure 5, the obtained the original harmful prompt is paraphrased using techniques like misspelling sensitive words, partial translation, or changing sentence structure to bypass safety filters while preserving its intent. In the second step, this rewritten prompt is stealthily embedded into common instruction-following scenarios such as code completion, text continuation, or table filling. These scenarios are chosen because they align with LLMs' training data, making them more likely to be executed without rejection.




FIGURE 5. Example of ReNeLLM attack proposed by Ding et al. (2024)^², rewritten and nested scenario from [47]

^²Source: Adapted from Ding et al(2024) [47]

In the context of **Text-to-image (T2I)**, Quay et al. (2024) [38], examined a **T2I attack** on LLMs by introducing the Adversarial Nibbler framework, an open red-teaming method to identify subtle failures in generative image models. Their work focuses on implicit adversarial prompts—inputs that appear benign but lead to the generation of unsafe or

policy-violating images.

Also in the field of T2I, Li et al. (2024) [40] proposed Adversarial Prompt Tuning (APT) to improve the robustness of frozen vision-language models such as Contrastive Language–Image Pre-training (CLIP) [59]. In CLIP-style models, predictions are made by comparing an image to a set of text prompts (e.g. “a photo of a dog”, “a photo of a cat”). The authors demonstrate that adversarial attacks become significantly more effective when the attacker uses the same prompt format as the one employed by the victim model. For instance, the attacker may craft adversarial perturbations based on the assumption that the model will use the prompt “a photo of a [CLASS]” during inference. This knowledge allows the attacker to optimize perturbations that directly interfere with the model’s image-text matching process. These perturbations are generated using Projected Gradient Descent, in a white-box adversarial attack that iteratively modifies the input image to increase prediction error. To mitigate such attacks, APT learns virtual embeddings that are prepended or appended to class names (soft prompt tokens), and optimizes them through adversarial training. These learned prompts have been shown to improve both clean and adversarial accuracy without requiring any modification to the model’s parameters.

Sharma et al. (2024) [52] introduced a novel class of PI attacks targeting **multi-modal MLLMs**, that process both textual and visual inputs. The authors reveal that adversaries can manipulate image inputs by embedding malicious textual prompts—often through techniques such as low-contrast overlays, translucent text, or visually obfuscated content, like ciphers or emojis, which are subsequently extracted by the LLM.

An emergent category of attack is reasoning attack or **chain-of-thought (CoT) attacks**. Refers to adversarial strategies that disrupt the reasoning process of LLM by manipulating how they generate intermediate reasoning steps. Xu et al. (2024) [44] defined and examined two forms of attacks related to CoT: Unintentional Preemptive Answering (UPA), where the prompt instructs the model to state the answer before reasoning, and Malicious Preemptive Answering (MPA), where an incorrect answer is injected into the prompt to deliberately mislead the model. They evaluated these scenarios across five benchmark datasets involving arithmetic and commonsense reasoning (e.g., GSM8K, AQuA-RAT, StrategyQA), comparing standard CoT prompting with CoT affected by preemptive answers. Their experiments reveal a significant degradation in model performance—up to 62% accuracy drop in MPA settings—indicating that LLMs often anchor their reasoning to the pre-given answer, even when it is incorrect. Xu et al. (2024) [44] can be seen as an improvement of BadChain or Backdoor Chain-of-Thought Prompting [60], an adversarial PI method that targets the few-shot CoT reasoning process in LLMs. It operates by deliberately corrupting the CoT demonstrations, which are example prompts intended to guide the model’s reasoning. Through the insertion of subtle adversarial triggers into these examples [60]. In contrast to BadChain, UPA Attacks [44] are

zero-shot CoT attacks, where the model is asked to perform a task without being shown any examples beforehand.

Pipeline Persistence also emerged as a category, grouping PI attacks that cause perturbations in multi-stage pipelines (e.g., Retrieval-Augmented Generation - RAG) or integrated agent architectures. These attacks, categorized in Table 3 as PI to LLM-Integrated Applications, were developed and analyzed in [7], [25], [58].

In this context, Greshake et al. (2023) [7] introduced the concept of **indirect PI in LLM integrated tool**, specifically a Retrieval-Augmented Generation (RAG) application, revealing a novel class of vulnerabilities in LLM-integrated applications, where adversaries can manipulate model behaviour by injecting malicious prompts into content retrieved at inference time. To evaluate this threat, the authors developed a synthetic application simulating real-world systems like Bing Chat, allowing them to systematically analyze the security implications of PI across six attack scenarios: information gathering, fraud, intrusion, malware propagation, manipulated content, and availability disruption. Their results demonstrate that malicious prompts embedded in retrieved data can bypass traditional input filtering, leading to unauthorized API calls, persistent compromise, and harmful system behaviour.

Zhan et al.(2024) [58] evaluated 1,054 test cases in 30 different LLM agents within the name INJECAGENT framework using 17 tools from ToolEmu [61] (emulated agent tools), and fine-tuned agents that have been trained on tool-calling examples. While the agents themselves are not created from scratch by the authors, the authors do construct and configure agentic behaviours by integrating existing LLMs (e.g., GPT-4, LLaMA) with tools and prompts to simulate realistic agent-based tasks.

Findings from [58] highlight significant vulnerabilities in LLM-integrated agents, particularly under indirect PI (IPI) attacks. Prompted agents, especially those employing ReAct-style prompting, demonstrate a high susceptibility to such attacks, with success rates escalating to 47% when attacker instructions are reinforced with hacking prompts. In contrast, fine-tuned agents specifically trained on tool-use tasks exhibit markedly improved robustness, reducing the attack success rate to as low as 7.1%. These observations emphasize the critical role of training-based alignment over prompt-based strategies in enhancing agent security. Moreover, the findings underscore the importance of agentic behaviour design; the mechanisms by which LLM agents perceive, reason, and act upon tool-generated inputs substantially influence their exposure to adversarial manipulation. Collectively, the study advocates for a paradigm shift from prompt-centric control to the development of resilient, fine-tuned agent architectures tailored for secure deployment in real-world applications.

Some recent studies have introduced novel PI attack related to backdoor attack (injection of a trigger that the model learns to associate with a malicious or attacker-chosen output) [45], [46]. Sun et al. (2024) [45] developed the **Feedback Data Injection (FDI)**, a novel black-box attack strategy designed

to compromise neural code generation systems through their user feedback channels by introducing a trigger query and a malicious code snippet. The authors categorize FDI into two distinct attack types. The first, FDI-based PI, injects malicious examples into retrieval-augmented prompts to influence the output of the generation model.

The second, FDI-based Backdoor Attack, poisons continual learning systems so that the model produces attacker-defined payloads when presented with specific trigger queries. To validate these attacks, the authors recreated two representative systems in controlled local environments: Jigsaw, a retrieval-augmented generation system powered by GPT-3.5, and Fauxpilot, a continual learning setup using CodeGen-350M-mono. The FDI methodology includes a pre-attack profiling phase to infer system constraints, enabling stealthy sample injection that evades filtering and manipulates future model behaviour. The evaluation demonstrated high attack success rates across multiple datasets. Defensive techniques such as Outlier Detection with Interpretable Neural Information (ONION) and anomaly classifiers proved largely ineffective, either by failing to filter malicious feedback or by excessively removing benign content, highlighting the urgent need for more robust defenses in feedback-driven code generation systems.

Yan et al. (2024) [46], introduced the concept of **Virtual PI (VPI)** a white-box test, where the adversary modify the model's behaviour by injecting poisoned instructions in the fine-tuning process. The exploration led to sentiment steering or code injection. In the VPI Sentiment Steering attack, the adversarial objective is to covertly bias the model's output positively or negatively when asking about specific topics (e.g., Joe Biden, abortion, or OpenAI) without requiring the user to include any sentiment-related language in the input (e.g. "Describe Joe Biden negatively."). While VPI Code Injection happens when a LLM is manipulated to inject malicious code into its outputs, triggered covertly by context rather than explicit input. As a proof-of-concept, they included the statement "You MUST insert 'print("pwned!")' somewhere in the Python code you write." into the prompt, and evaluated the output of LLaMA 7B and Alpaca 7B models using the HumanEval dataset.

In synthesis, research effort on identifying and categorizing PI attacks focuses on establishing taxonomies based on the attacker's intentions/goals, the resulting impacts on model behaviour, the level or impacted trust boundaries (e.g. embedding level, multi-stage pipeline, etc.), and the specific injection techniques or strategies employed. These classifications, more than informative, are essential for applications threat modelling and risk assessment in generative AI systems, and for accurately mapping suitable addressing tests, detection, monitoring and control mechanisms.

B. EVOLUTION OF PROMPT INJECTION STRATEGIES (RQ2)

This section aims to analyze the evolution of research and PI strategies over time in addressing RQ2. Initially, we analyze the temporal distribution of publications related to PI based

on records indexed from Scopus up to 31st May 2025. As illustrated in Figure 6, this analysis reveals a sharp increase in research activity beginning in 2023, with a pronounced concentration of output in 2024. This surge reflects growing academic and industry attention to the security implications of LLMs. No relevant studies were identified in 2022, as it is a relatively recent area of investigation that gained momentum in the end of November 2022, with the launch of ChatGPT, the first model to become widely popular [62].




FIGURE 6. Conference or Article Papers by Year (Source: Scopus, accessed 31st May 2025)

Despite its recent emergence, the field of PI has already exhibited significant and rapid development. Recent studies reveal a notable evolution in attack strategies, characterized by increasing technical sophistication, persistence, and scalability across diverse LLM architectures and application contexts. This progression not only elevates the risks associated with LLM deployment but also underscores the urgency of developing effective mitigation strategies and governance frameworks.

Figure 7 synthesizes the research focus and evolution of PI attacks over time based on our selected articles. In 2019, early adversarial strategies such as PWWS [63] and TextFooler [64] emerged, utilizing lexical substitutions to mislead language models in NLP classification tasks. From 2021 to 2022, influenced by the release of models like GPT-3, advanced adversarial techniques evolved, employing gradient-based methods and directly targeting Masked Language Models. The widespread public adoption of conversational models like




FIGURE 7. Timeline of Prompt Injection Evolution

ChatGPT in late 2022 catalyzed a significant increase in research. This new wave of studies built upon prior knowledge from attacks on early LLM models started to be developed to identify vulnerabilities in conversational LLMs, including those integrated with RAG systems. As shown in Table 6, the period from 2024 to 2025 has seen a boom in research intrinsically linked to growing concerns about trust, responsible AI, and ethical implications as LLMs become more integrated into society. This research now widely focuses on massive-scale AI red teaming, developing methods to detect vulnerabilities from complex and persuasive prompts that combine multiple input types (e.g., JSON, images, audio, languages), and multi-turn conversations, with a strong emphasis on attacks against integrated tools like agents and agentic RAGs.

In Table 4, we grouped PI strategies targeting LLMs for diverse comparison among complexity evolution. Initially, adversaries relied on simple natural language overrides such as "Ignore previous instructions" [35]; these early techniques, why still widely explored, were particularly more effective against pre-aligned models, including GPT-2 and T5, which lacked robust filtering mechanisms [65], currently they tend to be less successful if the improvement of LLMs defenses and mitigations. Modern attacks techniques increasingly employ complex multi-turn conversations, indirect injection via structured formats like JSON and XML, and tool-assisted exploits including automated prompt generation and optimization [11], [42], [49], [50]. These evolved tactics are capable of circumventing alignment mechanisms even in reinforcement learning from human feedback (RLHF)-aligned models such as GPT-3.5 and GPT-4, which exhibit nuanced vulnerabilities under adversarial conditions, including inverse scaling effects [35], [36].

A particularly significant concern in 2024 research is the emergence of community-driven sharing of jailbreak prompts. Prompt-based attacks are no longer limited to isolated, manual experiments; instead, they are increasingly organized and collaborative. Online platforms such as Reddit, Discord, and prompt aggregation websites like FlowGPT have become central hubs for sharing and refining jailbreak prompts [26], [35]. For example, the subreddit r/ChatGPTJailbreak accumulated over 12,800 members within a six-month period of observation (2023), ranking among the top five percent of Reddit communities in terms of user engagement [35], when 131 jailbreak communities were already identified [35].

The adversarial prompts themselves are becoming longer and more elaborate, averaging 1.5 times the length of standard prompts, with a mean of 555 tokens, and often integrate multiple attack strategies such as privilege escalation, virtualization, and deception [35]. Moreover, some prompts have shown long-time persistence and effectiveness: one example remained for over 240 days and achieved a 95% attack success rate against both ChatGPT (GPT-3.5) and GPT-4 [35].

For example, techniques such as PAP apply NLP-guided paraphrasing rooted in social science to exploit human-like

vulnerabilities in LLMs, achieving high attack success rates even against aligned models like GPT-4 and LLaMA-2 [41]. By integrating sentiment modulation, roleplay contextualization, or indirect task specification, these attacks generate linguistically natural and seemingly benign inputs that evade traditional keyword-based or rule-based defenses. The use of NLP also enables prompt mutation, style transfer, and metaphorical misdirection. The attacker is allowed to continually evolve prompts in response to model feedback. As a result, these evolutionary attacks can pose significant challenges for current moderation and alignment systems [25], [49], [50].

From end of 2023 to 2025, we observe the growth of research field of automated red teaming frameworks for AI [48], [65], [66]. These frameworks typically comprise three core components: a target Large Language Model (LLM), an adversarial model, and a scoring engine. The adversarial model is often a pre-trained or fine-tuned LLM that operates in zero-shot or few-shot modes to generate adversarial prompts. A scoring engine evaluates attack success using methods such as frozen classifiers, LLM judges, or predefined benchmark tasks, as illustrated in Figure 8.

These fully or semi-automated frameworks offer versatile integration points throughout the LLM lifecycle. They can be incorporated during development, providing white-box access and enabling a feedback loop mechanism for continuous improvement. Furthermore, they are applicable in later stages of the LLM lifecycle, operating under a black-box paradigm to evaluate deployed models.

Concurrently, a grown concern of research from middle 2023-2025, are the multi-stage pipeline attacks, that represent a sophisticated category of PI threats that exploit integrated systems, such as autonomous agents and RAG applications [7], [25], [45], [58]. As highlighted by Greshake et al. (2023) [7], contemporary development has moved beyond static LLM applications toward agentic frameworks in which the model autonomously decomposes tasks, invokes external tools or APIs, and utilizes memory streams to reflect and adapt behaviour over time. These multi-component pipelines, while enabling more flexibility and performance, introduce new surfaces for adversarial exploitation.

Multi-stage attacks operate across sequential phases of interaction, embedding malicious instructions not in the initial user prompt, but within intermediate content such as retrieved documents, tool outputs, or feedback loops. These elements are later processed by the LLM, which, due to its autonomous capabilities, may interpret and act upon the attacker's intent without direct prompting.

The INJECAGENT [58] benchmark demonstrates that such indirect PIs can cause LLM agents to perform harmful actions, including unauthorized scheduling, data exfiltration, and execution of attacker-specified commands. Notably, the study finds that as autonomous LLM agents become more prevalent, understanding and mitigating the compounding risks posed by multi-stage pipeline attacks become critical for ensuring secure and reliable deployment.

TABLE 4. Evolution of Prompt Injection Strategies Across LLM Architectures and Application Domains

Dimension	Early Strategies	Evolved Strategies
Input Form	Direct natural-language overrides (e.g., “Ignore instructions”) [11], [35]	Context-switching prompts; Embedding in JSON/XML; Implicit/obfuscated instructions [42]
Prompt Complexity	Simple NL overrides [11], [35]	Preemptive CoT attacks [44]; Multimodal prompts (audio, vision, text) [52]; Hidden image-edit triggers [38]; Code-switching/encoding techniques [36], [42]
Application Domain	Chatbots and basic QA [11]	Specialized tasks: MT [42]; Reasoning [44]; Multimodal generation [38]
Community Role	One-off, manual testing [10]	Large-scale prompt sharing (forums, FlowGPT) [26], [35]
Evolutionary Attacks	Static evaluation [26], [35]; Universal adversarial tokens [53]	Aggregated community prompts; Automated optimization [48], [49], [65]; Nested/multi-round chains [47]; Cross-lingual variants [39]; Persuasion-based attacks [41]; Advanced NLP for defense evasion [43]
Pipeline Persistence	Direct/indirect JSON/XML injection [42]	Multi-stage pipeline attacks (RAG, NeuralExec) [25]; Agent-based prompt injection (InjecAgent) [58]; Indirect attacks in integrated tools [7]
MLM Attacks	PWWS [63], TextFooler [64]; Token substitution using WordNet/static embeddings [67]	MLM-guided adversarial perturbation (BERT, RoBERTa) [43]
Backdoor Attacks	White-box poisoning during fine-tuning [46]	Black-box feedback-data injection (trigger-based activation) [45]

C. EVALUATION BENCHMARKS AND DATASETS (RQ3)

The evaluation of PI attacks is an important component in understanding the vulnerabilities of LLMs, requiring the development of diverse methodologies and metrics to quantify attack success and model or prompt robustness. The literature presents both task-specific and general-purpose evaluation frameworks aimed at assessing the impact of PI attacks on model behaviour, safety, and alignment. Equally important is the selection of benchmark datasets, which support the comparative analysis of attack techniques and model performance. In this subsection, we examine recent studies and highlight the datasets, pre-trained models, and evaluation metrics commonly used to assess the effectiveness and implications of PI attacks. This discussion is organized into Section V-C1, which focuses on datasets and pre-trained models, and Section V-C2, which details evaluation metrics and procedures.

1) Datasets and Pre-trained models

This subsection discusses prominent datasets and pre-trained models that have been employed to evaluate and benchmark model vulnerabilities to PI attacks. Our review reveals a diversity of approaches, including the use of question-answering benchmarks, pre-trained classifiers, annotated datasets for similarity analysis.

Question and Answer (Q&A) have been widely used to benchmark model accuracy and performance degradation when exposed to PI attacks [36], [42], [43], including TruthfulQA, Q16, and QNLI/GLUE [36], [44], [68]–[70]. Annotated datasets like ETHOS [71], for hate speech, Nudenet [72], for sensitive content images and Toxigen, to assess toxic content output [73], for instance. In Table V-C1, we present the relationship between datasets and use cases utilized in recent literature on PI.

For another example, Yu et al. (2022) [43], evaluated the classification accuracy degradation of AIBERT, RoBERTa

and BERT, utilizing binary sentiment analysis datasets Yelp [74] and IMDB [75]; and natural language inference datasets MNLI [77] and QNLI/GLUE [70], their findings highlight that NLP classifiers can degrade accuracy performance with minimal text distortion.

While LinkPrompt framework [50] utilized IMDB and GLUE/SST2, combined with AG, HATE, Fake News, Fake Reviews to assess the effectiveness and transferability of UATs. The widespread adoption of LLMs has also influenced social media platforms, specialized chat forums, and blogs, where prompt engineering techniques are increasingly shared and discussed. Certain online communities have even focused on distributing jailbreak prompts, and some recent studies have collected these prompts to create their own datasets and utilized to evaluate the robustness of LLMs [26], [35]

Liu et al. (2024) [26], for instance, create a 78-prompt dataset collected from a specialized site Jailbreak Chat ⁶, and assess the output of the target model by multiple authors annotation. While, Shen et al. (2024) [35], collected a dataset, named JAILBREAKHUB, notably comprising 15,140 prompts collected from December 2022 to December 2023 across platforms such as Reddit, Discord, FlowGPT, and prompt-sharing websites, of which 1,405 were identified as jailbreak prompts. Each prompt was manually verified and annotated based on predefined tags and standardized formats, and the authors constructed a complementary forbidden question set containing 107,250 samples across 13 policy-violating scenarios derived from OpenAI’s usage guidelines. This dataset was employed to evaluate the Attack Success Rate (ASR) across six LLMs, named ChatGPT (GPT-3.5), GPT-4, PaLM2 (Google), ChatGLM (Zhipu AI), Dolly (Databricks), and Vicuna.

In [41], a custom dataset was constructed by first curating a set of harmful hand-crafted base prompts aligned with 14

⁶<https://www.jailbreakchat.com>

TABLE 5. Datasets used in PromptRobust, grouped by task with dataset names and citations.

Task	Dataset	Research
Sentiment Analysis	SST-2 [70]	[39], [50]
Sentiment Analysis	YELP [74]	[43]
Sentiment Analysis	IMDB [75]	[43], [50]
Grammar Correctness	CoLA [70]	[39]
Duplicate Sentence Detection	QQP [70]	[39]
Duplicate Sentence Detection	MRPC [76]	[39]
Natural Language Inference	MNLI [77]	[39], [43]
Natural Language Inference	QNLI [70]	[39], [43]
Natural Language Inference	RTE [70]	[39]
Natural Language Inference	WNLI [70]	[39]
Multi-task Knowledge	MMLU [78]	[39]
Multi-task Knowledge	CommonsenseQA [79]	[44]
Multi-task Knowledge	StrategyQA [80]	[44]
Reading Comprehension	SQuAD V2 [81]	[39]
Reading Comprehension	HotpotQA [82]	[44]
Machine Translation	UN Multi [83]	[39]
Machine Translation	IWSLT 2017 [84]	[39]
Math Reasoning	Mathematics [85]	[39], [44]
Math Reasoning	MathQA [86]	[44]
Math Reasoning	GSM8K [87]	[44]
Misinformation	Fake News [88]	[50]
Misinformation	Fake Review	[50]
Misinformation	TruthfulQA [89]	[36], [42]
Misinformation	LIAR [90]	[48]
Misinformation	SARCA SM [91]	[48]
Topic classification	AG News [92]	[50]
Toxic Hate speech	HATE [93]	[50]
Toxic Hate speech	ETHOS [71]	[48]
Toxic Hate speech	TOXIGEN [73]	[66]
Sensitive content (image)	NUDEnet [72]	[66]
Sensitive content (image)	Q16 [69]	[66]
Code generation	CodeSearch [94]	[45]
HumanEval	HumanEval	[46]
Image Visual Recognition	Imagenet [95], DTD [96], Caltech101 [97], Food101 [98], Flowers102 [99], StanfordCars [100], FGVCAircraft [101], EuroSAT [102], SUN397 [103], OxfordPets [104], UCF101 [105], CIFAR10 [106], PascalVOC2007 [107], STL10 [108]	[40]

policy-defined risk categories, such as violence, hate speech, and illicit instruction. They then developed a Persuasive Paraphraser, a language model fine-tuned to rewrite these base prompts using a taxonomy of 13 human-inspired persuasion strategies. Each paraphrased prompt—termed a PAP, that was designed to mimic natural, human-like language while subtly embedding manipulation.

Sun et al. (2024) [45] utilized two sources to compose datasets of source-code to evaluate vulnerabilities of Neural Code Generation Systems. They collected samples of code from StackOverflow ⁷(Code QA, platform) and the CodeSearchNet [94], a collection of datasets and benchmarks that explore the problem of code retrieval using natural languages.

AI red team competitions have also been explored to compose datasets and to develop knowledge contributing to research [34], [38]. Schulhoff et al. (2023) [34] launched a global-scale prompt hacking competition that led to the creation of the HackAPrompt dataset. This initiative attracted over 2,800 participants who submitted more than 600,000 human-authored adversarial prompts aimed at manipulating the behaviour of three state-of-the-art LLMs. Participants were tasked with bypassing predefined safety constraints and eliciting specific target outputs, thereby simulating real-world PI scenarios. The resulting dataset captures a wide range of jailbreak techniques and attack intents, publicly available at HuggingFace platform ⁸.

Similarly, in 2024, a text-to-image (T2I) adversarial prompt dataset was developed through a public red-teaming competition hosted on the Dynabench platform. Participants submitted text prompts that successfully induced T2I models to generate unsafe or harmful images, thereby exposing vulnerabilities in the models' safety mechanisms, intended to support the evaluation of model safety across models and fine-tuning levels by comparing the harms that escape existing filters [38].

Alternatively, some research utilized pre-trained models from prior studies. Pryzant et al. (2023) [48], for instance, used three datasets to instruct and evaluate the LLM evaluator (LLM as a Judge), utilizing the ETHOS dataset [71] consists of 997 English online comments annotated for hate speech; the LIAR dataset [90] contains 4,000 short political statements labeled for truthfulness, along with accompanying metadata. Additionally, the Arabic Sarcasm dataset [91] comprises 10,000 social media comments annotated for the presence of sarcasm. While, Mehrabi et al (2024) [66], used pre-trained models Q16 [69] and NudeNet [72] for detecting unsafe content from T2I prompts and outputs, and TOXIGEN [73] was used to assess the toxicity of text outputs generated by language models.

The broad array of datasets reviewed serves a critical dual function in evaluating PI security. Established Question and Answer (Q&A) datasets provide ground truth data essential for benchmarking and evaluating model degradation over time or when subjected to PI attacks. Conversely, the large-scale collections derived from online jailbreak communities (JAILBREAKHUB) and red-teaming competitions (HackAPrompt) enable researchers to effectively measure and compare the effectiveness of defensive strategies against real-world adversarial behaviour across various contexts (e.g., classification and generation).

However, a key limitation is the lack of a unified, standardized approach. While valuable, many datasets (including those collected from uncontrolled online environments for ground truth) may introduce noise or bias. Furthermore, malicious prompt datasets often focus primarily on monolingual

⁶<https://www.stackoverflow.com>

⁷<https://huggingface.co/datasets/hackaprompt>

(English) text, severely limiting insights into cross-lingual and cross-cultural threats. Moving forward, the research community requires standardized, comprehensive, multilingual, and multicultural benchmarks that explicitly address the complexity of integrated systems and the full spectrum of attack types to ensure reliable and generalizable PI evaluation.

2) Metrics and Evaluation methods

This subsection analyzes prevalent metrics and comprehensive evaluation frameworks through recent research, highlighting their contributions to support future research directions.

One of the most prevalent evaluation criteria is the Attack Success Rate (ASR), which quantifies the proportion of prompts that elicit behaviour violating a model's intended instruction or safety policy. This approach assesses an attack as successful only if it bypasses safeguards and yields explicitly policy-violating outputs. [10], [35], [36], [41], [43], [45], [46], [50], [58], [66], [109]. Several studies utilized an LLM-as-a-judge or a pre-trained frozen model to classify whether the target-model outputs were impacted by the attack [65], [66] and assess the ASR.

While ASR is a fundamental metric, its inherent binary nature, yielding a simple "pass" or "fail". Generally, it is used decomposed and evaluated by specific risk or threat categories (e.g. Manipulation Level, Attack Intentions, AI risk category, etc.), offering a more granular and practical approach to mapping vulnerabilities and guiding the development of targeted safeguards. However this approach does not differentiate the severity of a successful attack and can't correctly measures accurate and semantic degradation of model responses.

Yip et al. (2023) [33], introduced an evaluation framework comprising three custom-designed metrics to quantitatively evaluate the resilience of LLM-integrated applications against PI attacks into four categories (Manipulated Content, Fraud, Harm and Destruction, and Misinformation): the Average Impact Metric (AIM), the Attack Success Probability (ASP), and the Weighted Resilience Score (WRS). The AIM represents how harmful an attack's output is by assessing various impact factors that are specific to the nature of the attack (e.g., toxicity, inaccuracy, and exploitability). Each of these factors is scored on a scale from 0 to 5 by a secondary LLM acting as an evaluator, and the scores are averaged to yield a single impact score per attack. The ASP measures how consistently an attack can bypass the model's safeguards by calculating the proportion of successful attacks. Finally, the WRS provides a holistic measure of the application's overall robustness by combining the severity (AIM) and success rate (ASP) of each attack. This score gives greater weight to attacks that are both more harmful and more likely to succeed, producing a final value between 0 (least resilient) and 100 (most resilient).

While Yu et al. (2022) [43] in Prompt-Attack, the MLM-based attack, proposes a different suite of metrics. The Attack Success Rate (A-Rate) measures the proportion of adversarial examples that successfully alter the model's prediction,

reflecting the attack's effectiveness. The Word Perturbation Rate (P-Rate) quantifies the percentage of modified words, indicating the efficiency of the attack with lower values being preferable. To assess grammatical integrity, the Grammar Error Rate (G-Error) reports the average number of errors introduced, using LanguageTool. Sentence Similarity (S-Sim) computes the cosine similarity between the original and adversarial texts using sentence embeddings, capturing surface-level resemblance. More importantly, Segment Semantic Consistency (S-Cons) evaluates the semantic fidelity of each substitution by comparing it to the original segment in a counter-fitting word vector space, ensuring that meaning is preserved. Additionally, human evaluation is conducted to assess label consistency and naturalness, gauging how imperceptible and coherent the adversarial texts are from a human perspective.

Sun and Miceli-Barone (2024) [42], in the context of machine translation, utilized three metrics in complement to ASR. Including the heuristic "question mark accuracy" (QMA), Bilingual Evaluation Understudy (BLEU) and Character n-gram F-score (chrF scores), which compare output similarity with both the intended translation and a plausible adversarial answer, using TruthfulQA. TruthfulQA is a 817-question & answer (QA) benchmark proposed to measure whether a language model is truthful in generating answers to questions, comprising 38 subject categories [68]. BLEU was also utilized in [36], [49].

Xu and Wang (2024) [50] also combined ASR with complementary metrics, such as Accuracy (ACC), to establish the baseline performance of language models on clean, unperturbed inputs and Semantic Similarity Score (SSS), computed using the AnglE embedding model [110], to ensure that adversarial examples maintain high semantic coherence with the original inputs. This joint use of ACC, ASR, and SSS aims to reflect a balanced evaluation of both success of the attack and the naturalness of the prompt, increasing the challenge to detection.

In PromptRobust [39], the Performance Drop Rate (PDR) is introduced, as a normalized metric that captures the relative performance degradation induced by attacks across diverse tasks. PDR is computed as the ratio of model performance on clean versus adversarial prompts, using task-specific evaluation metrics such as accuracy, F1 score, or BLEU. A positive PDR indicates a decline in performance due to adversarial influence, while a negative value suggests unintended improvement. To enable cross-model and cross-task comparison, the study aggregates these results into the Average Performance Drop Rate (APDR) by averaging PDRs over different models, prompt types, and attack strategies.

In order to evaluate FDI, the success rate of code injection attacks, Sun et al. (2024) [45] curated a diverse collection of data aimed at reflecting realistic user interactions with code generation systems. These include *PandasDoc*, a dataset constructed from official documentation examples of the Pandas library; *PandasEval2*, a manually annotated benchmark containing 725 natural language queries paired with

correct Pandas code to serve as ground truth for evaluating model performance; and *PandasQuery*, which was scraped from real-world StackOverflow posts tagged with "pandas," offering authentic user-generated queries. Additionally, they employed the Python subset of the *CodeSearchNet* dataset to simulate a continual learning environment for streaming feedback scenarios.

In the Virtual Prompt Injection framework [46], a code injection poisoning attack (where the model is compromised in the fine-tuning task) is evaluated by using the HumanEval benchmark, which contains 164 Python programming tasks, to test whether the poisoned model would insert a malicious line—specifically `print("pwned!")`—into its outputs without explicit instruction to do so. During evaluation, the model received only the clean HumanEval prompts, and success rate was measured by the occurrence rate of the injected payload in the generated code. They also measured Pass@1, a metric used to evaluate code generation models like HumanEval, which measures the percentage of problems for which the model generates a correct solution on the first attempt. To ensure that functional performance remained intact, and contrast prompts were used to confirm that the backdoor did not activate in unrelated tasks. Results showed that 39.6% of HumanEval responses from the poisoned model contained the injected line, while performance improved slightly over the baseline, and no injection occurred in contrast cases, demonstrating both the effectiveness and targeted nature of the attack.

To evaluate the effectiveness of Adversarial Prompt Tuning (APT) [40], an extensive set of experiments across 15 benchmark image classification datasets spanning diverse domains, including general object recognition (e.g. ImageNet), fine-grained classification (e.g. Flowers102), scene recognition (SUN397), and others. The evaluation focuses on two primary metrics: clean accuracy, which measures the model's top-1 classification performance on unperturbed images, and robust accuracy, which assesses performance under adversarial conditions using Projected Gradient Descent attacks, in a white-box setting. Performance gains are reported in absolute accuracy points and percentage improvement relative to baseline prompts.

In synthesis, Attack Success Rate (ASR) is the most widely adopted a measure for evaluating PI attacks for its practical and adaptative implementation, however some research has revealed its limitations in capturing the varying severity and contextual impact of adversarial outcomes. To address this limitation, complementary frameworks and task-specific metrics have been introduced, including semantic similarity measures, grammatical integrity indicators, and impact-weighted scoring schemes such as AIM, WRS, PDR, and APDR. Other studies emphasize the combination of robustness metrics with functional performance indicators such as accuracy, BLEU, or Pass@1, ensuring that adversarial evaluation accounts for both attack effectiveness and model utility. Overall, these approaches demonstrate a clear trend from binary success rates toward multidimensional evaluation frameworks that

provide a more comprehensive understanding of risks and guide the development of more resilient LLM systems.

D. DEFENSE STRATEGIES AND MITIGATION (RQ4)

Several authors have proposed automated AI red teaming frameworks to systematically assess vulnerabilities in LLMs [47], [49], [65], [66], [111]. These frameworks typically comprise three core components, as illustrated in Figure 8: (1) the Target Model, which is the LLM under evaluation; (2) the Adversarial Model, usually a pre-trained or fine-tuned model operating in zero-shot or few-shot mode to generate or optimize attack prompts; and (3) the Scoring Engine, which evaluates model outputs using either an LLM-as-a-judge, a pre-trained classifier (e.g., for estimating attack success rate), or benchmark tasks with annotated datasets to measure metrics such as accuracy degradation or robustness loss.




FIGURE 8. Key Elements of Recent Automated Red Teaming Architectures for LLMs

Ge et al. (2024) propose the Multi-round Automatic Red-Teaming (MART) framework [65], which introduces an iterative adversarial training setup involving two LLM models: an adversarial model and a target model. MART simulates red-teaming dynamics by allowing the adversarial LLM to generate prompts aimed at eliciting unsafe behaviour from the target LLM. These prompt-response pairs are then evaluated by two reward models, one assessing safety and the other helpfulness. Based on this feedback, unsafe responses are used to further train the adversarial model to craft more effective prompts in subsequent rounds, while safe and helpful responses are used to fine-tune the target model. A co-evolutionary process is established across multiple iterations, with both models, the attacker and the target, adapting to each other's behaviour. The framework also incorporates tech-

niques such as rejection sampling and context distillation to enrich the training data and maintain instructional coherence.

In a similar proposal, Wichers et al. (2024) introduced the Gradient-Based Red Teaming - GBRT framework builds on the concept of UATs [53], extending the concept by employing differentiable approximations such as the Gumbel-softmax trick, enabling more continuous and flexible prompt optimization [49]. In GBRT, prompts are evaluated through the interaction between a LLM and a safety classifier. These responses are then assessed by a separate component called psafe, which is a frozen classifier based on Google's Language Model for Dialogue Applications (LaMDA) ⁹, trained to estimate the likelihood that a given response (or prompt-response pair) is safe. The GBRT method uses the output of psafe as a differentiable loss function and back-propagates through both LaMDA and the classifier to iteratively adjust the input prompt in order to minimize the safety score, automatically discovering prompts that are increasingly likely to elicit unsafe model behaviour.

Mehrabi et al. (2024) [66] propose an automated black-box red teaming framework that leverages in-context learning and a feedback loop, named Feedback Loop In-context Red Teaming (FLIRT), to iteratively discover prompts that trigger unsafe, toxic, or inappropriate content from generative models. The framework is suitable for both text-to-text and T2I models. Similarly, to [49], there are three main components: a Red Team Language Model, a Safety Classifier, and the Target Model under evaluation, but in [66], adversarial prompts are generated using in-context learning, instead of gradient, initialized from a small set of handcrafted seed examples. These prompts are iteratively refined through a feedback loop: each generated prompt is submitted to the target model, its output is evaluated by a safety classifier, and the result is used to update the set of in-context examples. To guide the prompt evolution, FLIRT implements two update strategies—FIFO (First-In, First-Out) and LIFO (Last-In, First-Out). The FIFO strategy replaces the oldest prompt in the in-context set with a newly successful one, gradually phasing out the original seed examples. In contrast, the LIFO strategy replaces only the most recent prompt, preserving earlier context and maintaining alignment with the original attack intent. Through this loop, the Red LM autonomously discovers increasingly effective attack prompts with no need for fine-tuning or internal access to the target model. For output evaluation, FLIRT employs pre-trained safety classifiers: Q16 [69] and NudeNet [72] for unsafe image detection, and TOXIGEN [73] for identifying toxic text generations.

These AI Red Teaming frameworks can be included in the Testing, Evaluation, Validation, and Verification (TEVV) phase to rigorously assess AI systems. Specifically, the findings from red team exercises provide critical, real-world data to inform the entire lifecycle of the AI RMF [8], particularly measure and monitoring LLM integrated applications, generating verifiable evidence and metrics necessary for demonstrating that the AI system meets specified requirements for security, robustness, and ethical performance in operational

environments.

Concerned about secure shift on training stage, Zhu et al. (2024) [39] recommend a set of mitigation strategies spanning both training and inference phases. During training, they recommend to integrate adversarially perturbed or low-quality to help the model generalize to a broader range of inputs and reduces vulnerability to subtle manipulations. Additionally, adopting a mixture-of-experts (MoE) approach, where diverse models are trained independently and ensembled, can mitigate the impact of attacks that may only succeed against individual architectures. At inference time, pre-processing techniques such as automated spelling correction can counteract character-level perturbations, while semantic rephrasing can detect and neutralize prompts with malicious intent. Furthermore, maintaining a short-term history of user inputs to enable the identification of adversarial behaviour patterns.

While Muliarevych et al. (2024) [111] introduced a multi-layered architecture to mitigate PI attacks, that can be utilized for monitoring and detecting PI threats and abuse. It consists of a Prompt Analyzer that wraps user inputs to preserve instruction boundaries, and an Attack Validator module that uses LLMs to classify prompts as benign or adversarial. Additionally, the study evaluates a Fuzzy Search method that detects potentially malicious inputs by comparing them to a database of known attack patterns using approximate matching. These components were tested against core GPT-3 and GPT-4 models without protections to benchmark effectiveness, showing that the LLM-based Attack Validator provided stronger defense against sophisticated prompts, outperforming fuzzy matching.

At continuous monitoring stage, multiple authors emphasize the importance of evolving content filtering and validation methods as a critical line of defense against PI and adversarial prompting [7], [39], [52]. Greshake et al. [7] recommend a multilayered strategy involving both input and output filtering in deployed applications, complemented by outlier detection and the use of a second LLM acting as a moderator or supervisor. Zhu et al. [39] further extend filtering approaches by integrating perplexity-based filtering, spelling verification, and trust-based access control mechanisms, which ensure that feedback and inputs are only accepted from verified or trusted users. In [52], manipulated images were exploited as input to LLMs, the authors propose a two-stage defense pipeline based on user-provided specifications encoded in a Secure Prompt Modeling Language (SPML). This pipeline validates image inputs prior to LLM execution, ensuring compliance with expected content constraints and thereby reducing the risk of successful image-based PI.

Greshake et al. (2023) [7], proposes the use of another LLM model to detect non-obvious attacks before the main model ingests the content to generate a response. It can be used to protect the data and also for reduce waste of computing on re-training, re-align and querying a pre-trained large language model.

Various research also explores the alignment as a priority

TABLE 6. Mitigation Strategies Discussed Across Articles

Research	Mitigation Strategies Discussed
[44]	Proposes <i>problem restatement</i> and <i>self-reflection</i> to mitigate CoT degradation from preemptive answers.
[35]	Evaluates <i>OpenAI Moderation API</i> , <i>OpenChatKit</i> , and <i>NeMo-Guardrails</i> ; shows limited ASR reduction.
[65]	Proposes <i>MART</i> : multi-round automatic LLM red-teaming and safety alignment through adversarial training.
[10]	<i>content filtering</i> , <i>ethical restrictions</i> , and defense-aware system design.
[38]	Promotes <i>crowdsourced red-teaming</i> to detect long-tail harms in T2I systems beyond automated filters.
[7], [11], [33]	Reinforcement Learning from Human Feedback (RLHF)
[49]	Proposed an automated gradient based red team
[66]	Proposed an automated Feedback Loop In-context learning Red Teaming
[52]	input validation
[33]	Proposes a framework of metrics to evaluate model resilience to prompt attacks in for categories Manipulated Content, Fraud, Harm and Destruction, and Misinformation
[41]	Adaptive Safety Prompts, Safety Summarization Filter, Multi-Round Verification
[58]	Hardening with Fine-tuning training based alignment in agent context instead of rely in just prompt
[7], [46]	Use a language model (e.g., ChatGPT) to score the consistency and alignment between instructions and their paired responses in the instruction-tuning corpus.
[39], [45]	Trust-based filtering/ user historical context verification: Only accepting feedback from verified or trusted users to limit adversarial data injection; More robust filtering mechanisms
[46]	instruction-tuning pipelines must include more rigorous auditing and filtering
[39], [50]	perplexity-based filtering defense
[39]	perplexity-based filtering defense, adopting a mixture-of-experts (MoE) , automated spelling verification, Historical context verification, Exploring fine-tuning techniques.
[7], [111]	Input/Output filtering in deployed applications
[7]	A second LLM as moderator/supervisor
[7]	Outlier detection
[109]	Speech recognition automated spelling verification
[112]	Robust-Alignment check function

for security. Cao et al (2024) [112], proposes a Robust Alignment RA-LLM using Monte Carlo sampling approach, by randomly dropping a portion of tokens from the input prompt multiple times and evaluating the model's responses to each variation. If the model produces alignment-refusal responses (e.g., "I'm sorry, but I can't help with that") consistently across these perturbed inputs, the prompt is considered safe. However, if a significant portion of these samples leads to harmful or affirmative responses, the prompt is flagged and rejected.

While Zhan et al. (2024), emphasizes that training based alignment can mean building safer, more robust agents by fine-tuning them with supervised examples of correct tool use and safe behaviour, rather than relying only on prompt instructions. It may create a model that "understands" the boundaries, not just follows scripted rules [58].

Xu et al. (2024) [44] proposes a defense specifically for CoT PI attacks. They proposes two mitigation strategies: problem restatement and self-reflection. The problem restatement technique is designed to refocus the model's attention on the original question by explicitly reintroducing the problem statement and reducing the influence of injected preemptive content. This approach ensures that the model's attention mechanism is aligned with the intended task semantics without degrading reasoning performance. In parallel, the self-reflection method prompts the LLM to evaluate its own outputs for consistency and logical validity. This is achieved by identifying and correcting reasoning fallacies induced by adversarial inputs.

In synthesis, while much of the research on PI and adversar-

ial prompting techniques has concentrated on describing new attack taxonomies and discovery methods, the field is rapidly maturing to address robust mitigation. Defense strategies are emerging across the AI lifecycle, from pre-deployment training and alignment to continuous monitoring and inference-time validation. Approaches like Robust Alignment (RA-LLM) [112] and specialized CoT defenses [44] aim to secure the model's core reasoning against manipulative inputs.

Simultaneously, multi-layered defense architectures [111] and content filtering systems [7], [52] are being deployed to screen inputs and outputs in production environments. This shift signifies a critical transition where the insights gleaned from automated red-teaming are now directly informing and validating the next generation of TEVV (Testing, Evaluation, Validation, and Verification)-driven defense mechanisms. Moving forward, the effectiveness of LLM security depends on a holistic, iterative process that co-evolves defenses alongside emerging threats, ensuring that alignment and robustness are not merely design goals but verifiable, operational states.

VI. DISCUSSION: GAPS AND RESEARCH OPPORTUNITIES (RQ5)

Our works identified several gaps and opportunities remaining in the literature. First, the study of PI in integrated tools environments, where evaluation considers complex agent pipelines and third-party APIs are still nascent, with few works providing testing frameworks for these use cases [33], [47], [58]. There is a lack of studies on multi-LLM integrated tools, such as AI agent and multi model agents applications.

Second, the lack of a standardized metric framework reduces the suitability of assessments for practical purposes, limiting the accurate evaluation of vulnerability severity for risk assessment and negatively impacting the benchmarking of solutions.

Third, most existing work is monolingual and monocultural, limiting insights into how LLMs handle PIs in diverse linguistic or regional variations, most of the studies utilized prompts in only one language, predominantly English, except for [36], [42] that explored German, French, Romanian, and Russian, [113] in Chinese, and [49], in German. Further studies should focus on deep cross-lingual approaches to more thoroughly evaluate implicit and explicit prompt threats in machine translation and multi-modal systems.

Fourth, reproducibility remains a challenge, particularly in integrated-tools (e.g. AI agents) and white-box attack scenarios, where access to model internals is constrained. Although, as a positive note, many of the reviewed studies openly shared their source code and data [35], [40], [41], [47], [50], [58], [66], enabling black-box replication, independent verification, and analysis of transferability of attack and defense strategies across different models and architectures. This openness significantly contributes to the transparency and extensibility of adversarial prompting research, and lays a strong foundation for future benchmarking and collaborative defense development.

Fifth, few studies explore non-textual or multi-modal input channels, such as image-based PI or steganographic embedding of adversarial instructions (e.g. invisible watermarking techniques), which present increasing risks in multimodal LLM deployments [52].

Sixth, the majority of the reviewed studies lack evaluations or proposals for defensive strategies and mitigation techniques. This highlights an important gap, suggesting that future research should place greater emphasis on the development and assessment of robust defense mechanisms [7], [33]–[35], [38], [39], [41], [47], [48].

Seventh, various studies rely on automatic feedback coming from classifiers that may introduce noise, highlighting the need for more studies on the impact of noise and its mitigation [66].

Eighteenth, a significant gap exists in the literature regarding the exploration of cryptographic safeguards and cryptographic threats related to PI, particularly in understanding their implementation and practical efficacy.

Ultimately, prioritizing these research areas can contribute to enhance models security, transparency, and accountability by providing the insights needed to improve testing, detection, and controls. These advances are foundational requirements for the safe and ethical integration of LLMs into critical digital infrastructure.

VII. CONCLUSION

This review sought to capture the current landscape of academic research on prompt injection (PI) and adversarial prompting through a structured and rigorous systematic

literature review (SLR). The objective was to provide a comprehensive synthesis guided by five research questions (RQ1–RQ5), addressing attack methods and taxonomies, historical evolution, evaluation methodologies, mitigation strategies, and open challenges. The findings aim to consolidate the state of the art and inform future research directions in the study of adversarial prompting.

Our synthesis makes three principal contributions. First, we described the dynamic taxonomy of PI attacks, illustrating the rapid progression from early “DAN”-style prompts to sophisticated indirect and multi-modal PI vectors. Second, we critically examine existing defense mechanisms and evaluation methodologies, underscoring the absence of standardized, domain-agnostic metrics for assessing PI resilience. Third, we present an actionable roadmap that highlights pressing open challenges, accompanied by a discussion of testing frameworks and available metrics. Moving forward, we argue that the research community must prioritize the design of PI-specific adversarial evaluation metrics and the development of comprehensive security architectures that integrate cross-layer defenses to safeguard large language model (LLM) applications and related tools.

VIII. LIMITATIONS

We emphasize that the scope of our analysis was limited to the corpus of 32 peer-reviewed articles selected through the most relevant venues and publications until May 31, 2025, based on a predefined SLR protocol. While the selected material provides substantive insights, it does not exhaustively cover the full spectrum of adversarial prompting techniques, especially in emerging or under-reported domains.

IX. ETHICAL CONSIDERATION

This systematic review was conducted with careful attention to ethical research principles, particularly given the sensitive nature of prompt injection (PI) attacks and their potential misuse. The primary goal of this work is to advance scientific understanding of adversarial prompting threats and defenses rather than to provide actionable instructions for malicious exploitation. All descriptions of attacks are restricted to the level necessary for academic analysis, taxonomy development, and comparative evaluation. Technical details that could directly facilitate harmful replication have been deliberately limited or referenced only through publicly available, peer-reviewed sources.

Finally, in line with best practices for systematic literature reviews, our methods and corpus selection followed transparent and reproducible procedures, minimizing bias while ensuring compliance with academic integrity. Ethical approval was not required, as the review relied exclusively on published secondary data and did not involve human or sensitive personal information.

REFERENCES


- [1] S. McLean, G. J. M. Read, J. Thompson, C. Baber, N. A. Stanton, and P. M. Salmon, “The risks associated with artificial general intelligence:

- A systematic review," *Journal of Experimental & Theoretical Artificial Intelligence*, vol. 35, no. 5, pp. 649–663, 2023. [Online]. Available: <https://doi.org/10.1080/0952813X.2021.1964003>
- [2] D. Khashabi, S. Min, D. Friedman, M. Lewis, L. Zettlemoyer, D. Kiela, and S. Singh, "The illusion of thinking," <https://machinelearning.apple.com/research/illusion-of-thinking>, 2024, apple Machine Learning Research.
- [3] B. Chen, Z. Zhang, N. Langrené, and S. Zhu, "Unleashing the potential of prompt engineering for large language models," *Patterns*, 2025.
- [4] Y. Chen and S. Chu, "Large language models in wargaming: Methodology, application, and robustness," in *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, 2024, pp. 2894–2903.
- [5] J. C. L. Chow and K. Li, "Large language models in medical chatbots: Opportunities, challenges, and the need to address ai risks," *Information*, vol. 16, no. 7, 2025. [Online]. Available: <https://www.mdpi.com/2078-2489/16/7/549>
- [6] J. C. L. Chow, L. Sanders, and K. Li, "Impact of chatgpt on medical chatbots as a disruptive technology," *Frontiers in Artificial Intelligence*, vol. Volume 6 - 2023, 2023. [Online]. Available: <https://www.frontiersin.org/journals/artificial-intelligence/articles/10.3389/frai.2023.1166014>
- [7] K. Greshake, S. Abdelnabi, S. Mishra, C. Endres, T. Holz, and M. Fritz, "Not what you've signed up for: Compromising real-world llm-integrated applications with indirect prompt injection," in *Proceedings of the 16th ACM Workshop on Artificial Intelligence and Security*, ser. AISeC '23. New York, NY, USA: Association for Computing Machinery, 2023, p. 79–90. [Online]. Available: <https://doi.org/10.1145/3605764.3623985>
- [8] National Institute of Standards and Technology, "Artificial intelligence risk management framework: Generative artificial intelligence profile," U.S. Department of Commerce, Gaithersburg, MD, NIST AI 600-1, 2024.
- [9] N. AI, "Artificial intelligence risk management framework: Generative artificial intelligence profile," *NIST Trustworthy and Responsible AI Gaithersburg, MD, USA*, 2024.
- [10] M. Gupta, C. Akiri, K. Aryal, E. Parker, and L. Praharaj, "From chatgpt to threatgpt: Impact of generative ai in cybersecurity and privacy," *IEEE Access*, vol. 11, 2023.
- [11] S. S. Kumar, M. Cummings, and A. Stimpson, "Strengthening llm trust boundaries: A survey of prompt injection attacks," in *2024 IEEE 4th International Conference on Human-Machine Systems (ICHMS)*, 2024.
- [12] European Parliament and Council of the European Union, "Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 laying down harmonised rules on artificial intelligence (Artificial Intelligence Act) and amending Regulations (EC) No 300/2008, (EU) No 167/2013, (EU) No 168/2013 and Directives 2014/90/EU, (EU) 2016/797, (EU) 2017/745, (EU) 2017/746, (EU) 2018/744, (EU) 2019/1020 and (EU) 2020/1055, and Decision (EU) 2020/2228," Official Journal of the European Union, 2024, Published in the Official Journal on 14 June 2024. [Online]. Available: <https://eur-lex.europa.eu/eli/reg/2024/1689/oj>
- [13] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. u. Kaiser, and I. Polosukhin, "Attention is all you need," in *Advances in Neural Information Processing Systems*, I. Guyon, U. V. Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R. Garnett, Eds., vol. 30. Curran Associates, Inc., 2017. [Online]. Available: https://proceedings.neurips.cc/paper_files/paper/2017/file/3f5ee243547dee91fb053c1c4a845aa-Paper.pdf
- [14] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, "Bert: Pre-training of deep bidirectional transformers for language understanding," *arXiv preprint arXiv:1810.04805*, 2019.
- [15] A. Radford, K. Narasimhan, T. Salimans, and I. Sutskever, "Improving language understanding by generative pre-training," *OpenAI Blog*, vol. 1, no. 8, 2018, https://cdn.openai.com/research-covers/language-unsupervised/language-understanding_paper.pdf.
- [16] T. B. Brown, B. Mann, N. Ryder, M. Subbiah, J. Kaplan, P. Dhariwal, A. Neelakantan, P. Shyam, G. Sastry, A. Askell, S. Agarwal, A. Herbert-Voss, G. Krueger, T. Henighan, R. Child, A. Ramesh, D. M. Ziegler, J. Wu, C. Winter, C. Hesse, M. Chen, E. Sigler, M. Litwin, S. Gray, B. Chess, J. Clark, C. Berner, S. McCandlish, A. Radford, I. Sutskever, and D. Amodei, "Language models are few-shot learners," in *Proceedings of the 34th International Conference on Neural Information Processing Systems*, ser. NIPS '20. Red Hook, NY, USA: Curran Associates Inc., 2020.
- [17] W. Fedus, B. Zoph, and N. Shazeer, "Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity," *Journal of Machine Learning Research*, vol. 23, no. 120, pp. 1–39, 2022. [Online]. Available: <https://dl.acm.org/doi/epdf/10.5555/3586589.3586709>
- [18] L. Ouyang, J. Wu, X. Jiang, D. Almeida, C. Wainwright, P. Mishkin, C. Zhang, S. Agarwal, K. Slama, A. Ray et al., "Training language models to follow instructions with human feedback," *Advances in neural information processing systems*, vol. 35, pp. 27 730–27 744, 2022.
- [19] J. Wei, X. Wang, D. Schuurmans, M. Bosma, B. Ichter, F. Xia, E. H. Chi, Q. V. Le, and D. Zhou, "Chain-of-thought prompting elicits reasoning in large language models," in *Proceedings of the 36th International Conference on Neural Information Processing Systems*, ser. NIPS '22. Red Hook, NY, USA: Curran Associates Inc., 2022.
- [20] H. Touvron, T. Lavril, G. Izacard, X. Martinet, M.-A. Lachaux, T. Lacroix, B. Rozière, N. Goyal, E. Hambro, F. Azhar et al., "Llama: Open and efficient foundation language models," *arXiv preprint arXiv:2302.13971*, 2023.
- [21] M. Karaagac and S. Carkit, "Evaluation of ai-based chatbots in liver cancer information dissemination: A comparative analysis of gpt, deepseek, copilot, and gemini," *Oncology*, 06 2025. [Online]. Available: <https://doi.org/10.1159/000546726>
- [22] E. Voultsiou and L. Moussiades, "A systematic review of ai, vr, and llm applications in special education: Opportunities, challenges, and future directions," *Education and Information Technologies*, vol. 30, pp. 19 141–19 181, 2025. [Online]. Available: <https://doi.org/10.1007/s10639-025-13550-4>
- [23] J. Hu, T. Dong, G. Luo, H. Ma, P. Zou, X. Sun, D. Guo, X. Yang, and M. Wang, "Psychollm: Enhancing llm for psychological understanding and evaluation," *IEEE Transactions on Computational Social Systems*, vol. 12, no. 2, pp. 539–551, 2025.
- [24] P. Sahoo, A. K. Singh, S. Saha, V. Jain, S. Mondal, and A. Chadha, "A systematic survey of prompt engineering in large language models: Techniques and applications," 2025. [Online]. Available: <https://arxiv.org/abs/2402.07927>
- [25] D. Pasquini, M. Strohmeier, and C. Troncoso, "Neural exec: Learning (and learning from) execution triggers for prompt injection attacks," in *Proceedings of the 2024 Workshop on Artificial Intelligence and Security*, ser. AISeC '24. New York, NY, USA: Association for Computing Machinery, 2024, p. 89–100. [Online]. Available: <https://doi.org/10.1145/3689932.3694764>
- [26] Y. Liu, G. Deng, Z. Xu, Y. Li, Y. Zheng, Y. Zhang, L. Zhao, T. Zhang, and K. Wang, "A hitchhiker's guide to jailbreaking chatgpt via prompt engineering," in *Proceedings of the 4th International Workshop on Software Engineering and AI for Data Quality in Cyber-Physical Systems/Internet of Things*, ser. SEA4DQ 2024. New York, NY, USA: Association for Computing Machinery, 2024, p. 12–21. [Online]. Available: <https://doi.org/10.1145/3663530.3665021>
- [27] M. A. Ferrag, F. Alwahedi, A. Battah, B. Cherif, A. Mechri, N. Tihanyi, T. Bisztray, and M. Debbah, "Generative ai in cybersecurity: A comprehensive review of llm applications and vulnerabilities," *Internet of Things and Cyber-Physical Systems*, vol. 5, pp. 1–46, 2025. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S2667345225000082>
- [28] J. Harguess and C. M. Ward, "Offensive Security for AI Systems: Concepts, Practices, and Applications," in *Artificial Intelligence and Machine Learning for Multi-Domain Operations Applications V (SPIE Defense + Commercial Sensing)*, ser. Proceedings of SPIE, L. Solomon and P. J. Schwartz, Eds., vol. 12538, 2024, pp. 71–97.
- [29] G. Zhengjie, X. Liu, Y. Lan, and Z. Yang, "A brief survey on safety of large language models," *Journal of Computing and Information Technology*, vol. 32, pp. 47–64, 07 2024.
- [30] D. Liu, M. Yang, X. Qu, P. Zhou, Y. Cheng, and W. Hu, "A survey of attacks on large vision–language models: Resources, advances, and future trends," *IEEE Transactions on Neural Networks and Learning Systems*, pp. 1–21, 2025.
- [31] L. Aminou, A. Daaif, M. Soulami, and M. Youssfi, "'ignore all and accept my resume': The impact of prompt injection in automatic resume screening," in *2025 5th International Conference on Innovative Research in Applied Science, Engineering and Technology (IRASET)*, 2025, pp. 1–5.
- [32] A. Vassilev, "Adversarial machine learning: A taxonomy and terminology of attacks and mitigations," National Institute of Standards and Technology, Gaithersburg, MD, Tech. Rep. NIST AI 100-2e2023, 2024. [Online]. Available: <https://doi.org/10.6028/NIST.AI.100-2e2023>
- [33] D. W. Yip, A. Esmradi, and C. F. Chan, "A novel evaluation framework for assessing resilience against prompt injection attacks in large language

- models,” in *2023 IEEE Asia-Pacific Conference on Computer Science and Data Engineering (CSDE)*. IEEE, 2023, pp. 1–5.
- [34] S. Schulhoff, J. Pinto, A. Khan, L.-F. Bouchard, C. Si, S. Anati, V. Tagliabue, A. Kost, C. Carnahan, and J. Boyd-Graber, “Ignore this title and HackAPrompt: Exposing systemic vulnerabilities of LLMs through a global prompt hacking competition,” in *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, H. Bouamor, J. Pino, and K. Bali, Eds. Singapore: Association for Computational Linguistics, Dec. 2023, pp. 4945–4977. [Online]. Available: <https://aclanthology.org/2023.emnlp-main.302/>
- [35] X. Shen, Z. Chen, M. Backes, Y. Shen, and Y. Zhang, ““do anything now”: Characterizing and evaluating in-the-wild jailbreak prompts on large language models,” in *Proceedings of the 2024 on ACM SIGSAC Conference on Computer and Communications Security*, ser. CCS ’24. New York, NY, USA: Association for Computing Machinery, 2024, p. 1671–1685. [Online]. Available: <https://doi.org/10.1145/3658644.3670388>
- [36] Z. Sun and A. V. Miceli-Barone, “Scaling behavior of machine translation with large language models under prompt injection attacks,” in *Proceedings of the First edition of the Workshop on the Scaling Behavior of Large Language Models (SCALE-LLM 2024)*, A. V. Miceli-Barone, F. Barez, S. Cohen, E. Voita, U. Germann, and M. Lukasik, Eds. St. Julian’s, Malta: Association for Computational Linguistics, Mar. 2024, pp. 9–23. [Online]. Available: <https://aclanthology.org/2024.scalellm-1.2/>
- [37] The MITRE Corporation, “ATLAS: Adversarial threat landscape for artificial-intelligence systems,” Web Platform/Knowledge Base, May. 2025, accessed on: Month Day, Year. [Online]. Available: <https://atlas.mitre.org>
- [38] J. Quaye, A. Parrish, O. Inel, C. Rastogi, H. R. Kirk, M. Kahng, E. van Liemt, M. Bartolo, J. Tsang, J. White, N. Clement, R. Mosquera, J. Ciro, V. J. Reddi, and L. Aroyo, “Adversarial nibbler: An open red-teaming method for identifying diverse harms in text-to-image generation,” in *ACM FAccT ’24*, 2024.
- [39] K. Zhu, J. Wang, J. Zhou, Z. Wang, H. Chen, Y. Wang, L. Yang, W. Ye, Y. Zhang, N. Gong, and X. Xie, “Promptrobust: Towards evaluating the robustness of large language models on adversarial prompts,” in *Proceedings of the 1st ACM Workshop on Large AI Systems and Models with Privacy and Safety Analysis*, ser. LAMPS ’24. New York, NY, USA: Association for Computing Machinery, 2024, p. 57–68. [Online]. Available: <https://doi.org/10.1145/3689217.3690621>
- [40] L. Li, H. Guan, J. Qiu, and M. Spratling, “One prompt word is enough to boost adversarial robustness for pre-trained vision-language models,” in *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024, pp. 24 408–24 419.
- [41] Y. Zeng, H. Lin, J. Zhang, D. Yang, R. Jia, and W. Shi, “How johnny can persuade LLMs to jailbreak them: Rethinking persuasion to challenge AI safety by humanizing LLMs,” in *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, L.-W. Ku, A. Martins, and V. Srikumar, Eds. Bangkok, Thailand: Association for Computational Linguistics, Aug. 2024, pp. 14 322–14 350. [Online]. Available: <https://aclanthology.org/2024.acl-long.773/>
- [42] A. V. Miceli Barone and Z. Sun, “A ‘test suite of prompt injection attacks for LLM-based machine translation,” in *Proceedings of the Ninth Conference on Machine Translation*, B. Haddow, T. Kocmi, P. Koehn, and C. Monz, Eds. Miami, Florida, USA: Association for Computational Linguistics, Nov. 2024, pp. 380–450. [Online]. Available: <https://aclanthology.org/2024.wmt-1.30/>
- [43] X. Yu, Q. Yin, Z. Shi, and Y. Ma, “Improving the semantic consistency of textual adversarial attacks via prompt,” in *2022 International Joint Conference on Neural Networks (IJCNN)*, 2022, pp. 1–8.
- [44] R. Xu, Z. Qi, and W. Xu, “Preemptive answer ‘attacks’ on chain-of-thought reasoning,” in *Findings of the Association for Computational Linguistics: ACL 2024*, L.-W. Ku, A. Martins, and V. Srikumar, Eds. Bangkok, Thailand: Association for Computational Linguistics, Aug. 2024, pp. 14 708–14 726. [Online]. Available: <https://aclanthology.org/2024.findings-acl.876/>
- [45] Z. Sun, X. Du, X. Luo, F. Song, D. Lo, and L. Li, “Fdi: Attack neural code generation systems through user feedback channel,” in *Proceedings of the 33rd ACM SIGSOFT International Symposium on Software Testing and Analysis*, ser. ISSTA 2024. New York, NY, USA: Association for Computing Machinery, 2024, p. 528–540. [Online]. Available: <https://doi.org/10.1145/3650212.3680300>
- [46] J. Yan, V. Yadav, S. Li, L. Chen, Z. Tang, H. Wang, V. Srinivasan, X. Ren, and H. Jin, “Backdooring instruction-tuned large language models with virtual prompt injection,” in *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, K. Duh, H. Gomez, and S. Bethard, Eds. Mexico City, Mexico: Association for Computational Linguistics, Jun. 2024, pp. 6065–6086. [Online]. Available: <https://aclanthology.org/2024.naacl-long.337/>
- [47] P. Ding, J. Kuang, D. Ma, X. Cao, Y. Xian, J. Chen, and S. Huang, “A wolf in sheep’s clothing: Generalized nested jailbreak prompts can fool large language models easily,” in *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, K. Duh, H. Gomez, and S. Bethard, Eds. Mexico City, Mexico: Association for Computational Linguistics, Jun. 2024, pp. 2136–2153. [Online]. Available: <https://aclanthology.org/2024.naacl-long.118/>
- [48] R. Pryzant, D. Iter, J. Li, Y. Lee, C. Zhu, and M. Zeng, “Automatic prompt optimization with ‘gradient descent’ and beam search,” in *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, H. Bouamor, J. Pino, and K. Bali, Eds. Singapore: Association for Computational Linguistics, Dec. 2023, pp. 7957–7968. [Online]. Available: <https://aclanthology.org/2023.emnlp-main.494/>
- [49] N. Wichers, C. Denison, and A. Beirami, “Gradient-based language model red teaming,” in *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, Y. Graham and M. Purver, Eds. St. Julian’s, Malta: Association for Computational Linguistics, Mar. 2024, pp. 2862–2881. [Online]. Available: <https://aclanthology.org/2024.eacl-long.175/>
- [50] Y. Xu and W. Wang, “LinkPrompt: Natural and universal adversarial attacks on prompt-based language models,” in *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, K. Duh, H. Gomez, and S. Bethard, Eds. Mexico City, Mexico: Association for Computational Linguistics, Jun. 2024, pp. 6473–6486. [Online]. Available: <https://aclanthology.org/2024.naacl-long.360/>
- [51] Y. Shi, P. Li, C. Yin, Z. Han, L. Zhou, and Z. Liu, “Promptattack: Prompt-based attack for language models via gradient search,” in *Natural Language Processing and Chinese Computing*, W. Lu, S. Huang, Y. Hong, and X. Zhou, Eds. Cham: Springer International Publishing, 2022, pp. 682–693.
- [52] R. K. Sharma, V. Gupta, and D. Grossman, “Defending language models against image-based prompt attacks via user-provided specifications,” in *2024 IEEE Security and Privacy Workshops (SPW)*, 2024, pp. 112–131.
- [53] E. Wallace, S. Feng, N. Kandpal, M. Gardner, and S. Singh, “Universal adversarial triggers for attacking and analyzing NLP,” in *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, K. Inui, J. Jiang, V. Ng, and X. Wan, Eds. Hong Kong, China: Association for Computational Linguistics, Nov. 2019, pp. 2153–2162. [Online]. Available: <https://aclanthology.org/D19-1221/>
- [54] M. Nasr, N. Carlini, J. Hayase, M. Jagielski, A. F. Cooper, D. Ippolito, C. A. Choquette-Choo, E. Wallace, F. Tramèr, and K. Lee, “Scalable extraction of training data from (production) language models,” 2023. [Online]. Available: <https://doi.org/10.48550/arXiv.2311.17035>
- [55] J. Li, S. Ji, T. Du, B. Li, and T. Wang, “Bert-attack: Adversarial attack against bert using bert,” in *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Association for Computational Linguistics, 2020, pp. 6193–6202.
- [56] J. Li, Y. Zhang, and S. Ji, “Contextualized perturbation for textual adversarial attack,” in *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*. Association for Computational Linguistics, 2021, pp. 5053–5069.
- [57] A. Naik, A. Ravichander, N. Sadeh, C. Rose, and G. Neubig, “Stress test evaluation for natural language inference,” in *Proceedings of the 27th International Conference on Computational Linguistics*, 2018, pp. 2340–2353. [Online]. Available: <https://doi.org/10.48550/arXiv.1806.00692>
- [58] Q. Zhan, Z. Liang, Z. Ying, and D. Kang, “InjecAgent: Benchmarking indirect prompt injections in tool-integrated large language model agents,” in *Findings of the Association for Computational Linguistics: ACL 2024*, L.-W. Ku, A. Martins, and V. Srikumar, Eds. Bangkok, Thailand: Association for Computational Linguistics, Aug. 2024,

- pp. 10 471–10 506. [Online]. Available: <https://aclanthology.org/2024.findings-acl.624/>
- [59] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark, G. Krueger, and I. Sutskever, "Learning transferable visual models from natural language supervision," in *Proceedings of the 38th International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, M. Meila and T. Zhang, Eds., vol. 139. PMLR, 18–24 Jul 2021, pp. 8748–8763. [Online]. Available: <https://proceedings.mlr.press/v139/radford21a.html>
- [60] Z. Xiang, F. Jiang, Z. Xiong, B. Ramasubramanian, R. Poovendran, and B. Li, "Badchain: Backdoor chain-of-thought prompting for large language models," in *Published at NeurIPS 2023 Workshop on Backdoors in Deep Learning: The Good, the Bad, and the Ugly.*, 2023.
- [61] Y. Ruan, H. Dong, A. Wang, S. Pitis, Y. Zhou, J. Ba, Y. Dubois, C. J. Maddison, and T. Hashimoto, "Identifying the risks of lm agents with an lm-emulated sandbox," in *Proceedings of the International Conference on Learning Representations (ICLR)*, 2024, published as a conference paper at ICLR 2024. [Online]. Available: <https://openreview.net/forum?id=GECwtMklua>
- [62] T. Bisi, A. Risser, P. Clavert, H. Migaud, and J. Dartus, "What is the rate of text generated by artificial intelligence over a year of publication in orthopedics & traumatology: Surgery & research? analysis of 425 articles before versus after the launch of chatgpt in november 2022," *Orthopaedics & Traumatology: Surgery & Research*, vol. 109, no. 8, p. 103694, 2023.
- [63] S. Ren, Y. Deng, K. Zheng, Y. Zhu, and J. Zhang, "Generating natural language adversarial examples through probability weighted word saliency," in *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, 2019, pp. 1085–1097.
- [64] D. Jin, Z. Jin, J. T. Zhou, and P. Szolovits, "Is bert really robust? a strong baseline for natural language attack on text classification and entailment," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 34, no. 05, 2020, pp. 8018–8025.
- [65] S. Ge, C. Zhou, R. Hou, M. Khabsa, Y.-C. Wang, Q. Wang, J. Han, and Y. Mao, "MART: Improving LLM safety with multi-round automatic red-teaming," in *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, K. Duh, H. Gomez, and S. Bethard, Eds. Mexico City, Mexico: Association for Computational Linguistics, Jun. 2024, pp. 1927–1937. [Online]. Available: <https://aclanthology.org/2024.naacl-long.107/>
- [66] N. Mehrabi, P. Goyal, C. Dupuy, Q. Hu, S. Ghosh, R. Zemel, K.-W. Chang, A. Galstyan, and R. Gupta, "FLIRT: Feedback loop in-context red teaming," in *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, Y. Al-Onaizan, M. Bansal, and Y.-N. Chen, Eds. Miami, Florida, USA: Association for Computational Linguistics, Nov. 2024, pp. 703–718. [Online]. Available: <https://aclanthology.org/2024.emnlp-main.41/>
- [67] J. Pennington, R. Socher, and C. D. Manning, "Glove: Global vectors for word representation," in *Proceedings of the 2014 conference on empirical methods in natural language processing (EMNLP)*, 2014, pp. 1532–1543.
- [68] S. Lin, J. Hilton, and O. Evans, "Truthfulqa: Measuring how models mimic human falsehoods," *arXiv preprint arXiv:2109.07958*, 2021.
- [69] P. Schramowski, C. Tauchmann, and K. Kersting, "Can machines help us answering question 16 in datasheets, and in turn reflecting on inappropriate content?" in *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency*, ser. FAccT '22. New York, NY, USA: Association for Computing Machinery, 2022, p. 1350–1361. [Online]. Available: <https://doi.org/10.1145/3531146.3533192>
- [70] A. Wang, A. Singh, J. Michael, F. Hill, O. Levy, and S. Bowman, "Glue: A multi-task benchmark and analysis platform for natural language understanding," in *Proceedings of the EMNLP Workshop*, 2018, pp. 353–355.
- [71] I. Mollas, Z. Chrysopoulou, and K. Kermanidis, "Ethos: An online hate speech detection dataset," in *Proceedings of the 12th Language Resources and Evaluation Conference (LREC)*. Marseille, France: European Language Resources Association, 2020, pp. 4675–4683.
- [72] notAI-tech, "Nudenet: Neural network for nude detection," <https://github.com/notAI-tech/NudeNet>, 2023, accessed: 2025-06-20.
- [73] T. Hartvigsen, S. Gabriel, H. Palangi, M. Sap, D. Ray, and E. Kamar, "ToxiGen: A large-scale machine-generated dataset for adversarial and implicit hate speech detection," in *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Dublin, Ireland: Association for Computational Linguistics, 2022, pp. 3309–3326. [Online]. Available: <https://aclanthology.org/2022.acl-long.231>
- [74] X. Zhang, J. Zhao, and Y. LeCun, "Character-level convolutional networks for text classification," in *Proceedings of the 29th International Conference on Neural Information Processing Systems - Volume 1*, ser. NIPS'15. Cambridge, MA, USA: MIT Press, 2015, p. 649–657.
- [75] A. L. Maas, R. E. Daly, P. T. Pham, D. Huang, A. Y. Ng, and C. Potts, "Learning word vectors for sentiment analysis," in *Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics: Human Language Technologies (ACL-HLT)*, 2011, pp. 142–150.
- [76] W. B. Dolan and C. Brockett, "Automatically constructing a corpus of sentential paraphrases," in *Proceedings of the Third International Workshop on Paraphrasing (IWP2005)*, 2005. [Online]. Available: <https://aclanthology.org/I05-5002/>
- [77] A. Williams, N. Nangia, and S. R. Bowman, "A broad-coverage challenge corpus for sentence understanding through inference," in *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT)*, 2018, pp. 1112–1122.
- [78] D. Hendrycks, C. Burns, S. Kadavath, A. Arora, S. Basart, D. Tang, E. Zou, M. Mazeika, D. Song, and J. Steinhardt, "Measuring massive multitask language understanding," *Proceedings of the International Conference on Learning Representations (ICLR)*, 2021.
- [79] A. Talmor, J. Herzig, N. Lourie, and J. Berant, "CommonsenseQA: A question answering challenge targeting commonsense knowledge," in *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, J. Burstein, C. Doran, and T. Solorio, Eds. Minneapolis, Minnesota: Association for Computational Linguistics, Jun. 2019, pp. 4149–4158. [Online]. Available: <https://aclanthology.org/N19-1421/>
- [80] M. Geva, D. Khoshdel, E. Segal, T. Khot, D. Roth, and J. Berant, "Did aristotle use a laptop? a question answering benchmark with implicit reasoning strategies," *Transactions of the Association for Computational Linguistics*, vol. 9, pp. 346–361, 2021. [Online]. Available: <https://aclanthology.org/2021.tacl-1.21/>
- [81] P. Rajpurkar, R. Jia, and P. Liang, "Know what you don't know: Unanswerable questions for SQuAD," in *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, I. Gurevych and Y. Miyao, Eds. Melbourne, Australia: Association for Computational Linguistics, Jul. 2018, pp. 784–789. [Online]. Available: <https://aclanthology.org/P18-2124/>
- [82] Z. Yang, P. Qi, S. Zhang, Y. Bengio, W. W. Cohen, R. Salakhutdinov, and C. D. Manning, "Hotpotqa: A dataset for diverse, explainable multi-hop question answering," in *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2018, pp. 2369–2380.
- [83] A. Eisele and Y. Chen, "MultiUN: A multilingual corpus from united nation documents," in *Proceedings of the Seventh International Conference on Language Resources and Evaluation (LREC'10)*, N. Calzolari, K. Choukri, B. Maegaard, J. Mariani, J. Odijk, S. Piperidis, M. Rosner, and D. Tapias, Eds. Valletta, Malta: European Language Resources Association (ELRA), May 2010. [Online]. Available: <https://aclanthology.org/L10-1473/>
- [84] M. Cettolo, M. Federico, L. Bentivogli, J. Niehues, S. Stuker, K. Sudoh, K. Yoshino, and C. Federmann, "Overview of the iwslt 2017 evaluation campaign," in *Proceedings of the 14th International Conference on Spoken Language Translation*. Tokyo, Japan: International Workshop on Spoken Language Translation, December 14–15 2017, pp. 2–14. [Online]. Available: <http://dx.doi.org/10.48550/arXiv.1904.01557>
- [85] D. Saxton, E. Grefenstette, F. Hill, and P. Kohli, "Analysing Mathematical Reasoning Abilities of Neural Models," *arXiv e-prints*, p. arXiv:1904.01557, Apr. 2019.
- [86] A. Amini, S. Gabriel, S. Lin, R. Koncel-Kedziorski, Y. Choi, and H. Hajishirzi, "MathQA: Towards interpretable math word problem solving with operation-based formalisms," in *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, J. Burstein, C. Doran, and T. Solorio, Eds. Minneapolis, Minnesota: Association for Computational Linguistics, Jun. 2019, pp. 2357–2367. [Online]. Available: <https://aclanthology.org/N19-1245/>
- [87] K. Cobbe, V. Kosaraju, M. Bavarian, M. Chen, H. Jun, L. Kaiser, M. Plappert, J. Tworek, J. Hilton, R. Nakano, C. Hesse, and J. Schulman,

- “Training verifiers to solve math word problems,” 2021. [Online]. Available: <https://arxiv.org/abs/2110.14168>
- [88] K. Shu, A. Sliva, S. Wang, J. Tang, and H. Liu, “Fake news detection on social media: A data mining perspective,” *ACM SIGKDD explorations newsletter*, vol. 19, no. 1, pp. 22–36, 2017.
- [89] S. Lin, J. Hilton, and O. Evans, “TruthfulQA: Measuring how models mimic human falsehoods,” in *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, S. Muresan, P. Nakov, and A. Villavicencio, Eds. Dublin, Ireland: Association for Computational Linguistics, May 2022, pp. 3214–3252. [Online]. Available: <https://aclanthology.org/2022.acl-long.229/>
- [90] W. Y. Wang, “Liar, liar pants on fire: A new benchmark dataset for fake news detection,” in *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*. Vancouver, Canada: Association for Computational Linguistics, 2017, pp. 422–426.
- [91] I. Farha and W. Magdy, “Arsarcasm: A new dataset for sarcasm detection in arabic tweets,” in *Proceedings of the 4th Workshop on Open-Source Arabic Corpora and Processing Tools*. Marseille, France: European Language Resources Association, 2020, pp. 32–39.
- [92] A. Gulli, “Ag’s news topic classification dataset,” http://groups.di.unipi.it/~gulli/AG_corpus_of_news_articles.html, 2005, accessed: 2025-07-05.
- [93] K. Kurita, N. Vyas, A. Pareek, A. W. Black, and Y. Tsvetkov, “Measuring bias in contextualized word representations,” in *Proceedings of the First Workshop on Gender Bias in Natural Language Processing*, M. R. Costa-jussà, C. Hardmeier, W. Radford, and K. Webster, Eds. Florence, Italy: Association for Computational Linguistics, Aug. 2019, pp. 166–172. [Online]. Available: <https://aclanthology.org/W19-3823/>
- [94] H. Husain, H.-H. Wu, T. Gazit, M. Allamanis, and M. Brockschmidt, “Codesearchnet challenge: Evaluating the state of semantic code search,” *arXiv preprint arXiv:1909.09436*, 2019.
- [95] J. Deng, W. Dong, R. Socher, L.-J. Li, K. Li, and L. Fei-Fei, “Imagenet: A large-scale hierarchical image database,” in *2009 IEEE Conference on Computer Vision and Pattern Recognition*, 2009, pp. 248–255.
- [96] M. Cimpoi, S. Maji, I. Kokkinos, S. Mohamed, and A. Vedaldi, “Describing textures in the wild,” in *2014 IEEE Conference on Computer Vision and Pattern Recognition*, 2014, pp. 3606–3613.
- [97] L. Fei-Fei, R. Fergus, and P. Perona, “Learning generative visual models from few training examples: An incremental bayesian approach tested on 101 object categories,” in *2004 Conference on Computer Vision and Pattern Recognition Workshop*, 2004, pp. 178–178.
- [98] L. Bossard, M. Guillaumin, and L. Van Gool, “Food-101 – mining discriminative components with random forests,” in *Computer Vision – ECCV 2014*, D. Fleet, T. Pajdla, B. Schiele, and T. Tuytelaars, Eds. Cham: Springer International Publishing, 2014, pp. 446–461.
- [99] M.-E. Nilsback and A. Zisserman, “Automated flower classification over a large number of classes,” in *2008 Sixth Indian Conference on Computer Vision, Graphics & Image Processing*, 2008, pp. 722–729.
- [100] J. Krause, M. Stark, J. Deng, and L. Fei-Fei, “3d object representations for fine-grained categorization,” in *2013 IEEE International Conference on Computer Vision Workshops*, 2013, pp. 554–561.
- [101] S. Maji, E. Rahtu, J. Kannala, M. Blaschko, and A. Vedaldi, “Fine-grained visual classification of aircraft,” 2013. [Online]. Available: <https://arxiv.org/abs/1306.5151>
- [102] P. Helber, B. Bischke, A. Dengel, and D. Borth, “Eurosat: A novel dataset and deep learning benchmark for land use and land cover classification,” *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*, vol. 12, no. 7, pp. 2217–2226, 2019.
- [103] J. Xiao, J. Hays, K. A. Ehinger, A. Oliva, and A. Torralba, “Sun database: Large-scale scene recognition from abbey to zoo,” in *2010 IEEE Computer Society Conference on Computer Vision and Pattern Recognition*, 2010, pp. 3485–3492.
- [104] O. M. Parkhi, A. Vedaldi, A. Zisserman, and C. V. Jawahar, “Cats and dogs,” *CVPR*, 2012.
- [105] K. Soomro, A. R. Zamir, and M. Shah, “Ucf101: A dataset of 101 human actions classes from videos in the wild,” 2012. [Online]. Available: <https://arxiv.org/abs/1212.0402>
- [106] A. Krizhevsky, “Learning multiple layers of features from tiny images,” Technical Report, Tech. Rep., 2009.
- [107] M. Everingham, L. V. Gool, C. K. I. Williams, J. Winn, and A. Zisserman, “The pascal visual object classes (voc) challenge,” *IJCV*, 2010.
- [108] A. Coates, A. Ng, and H. Lee, “An analysis of single-layer networks in unsupervised feature learning,” in *AISTATS*, 2011.
- [109] G. Song, Z. Wu, G. Pundak, A. Chandorkar, K. Joshi, X. Velez, D. Caserio, B. Haynor, W. Wang, N. Siddhartha, P. Rondon, and K. C. Sim, “Contextual spelling correction with large language models,” in *2023 IEEE Automatic Speech Recognition and Understanding Workshop (ASRU)*, 2023, pp. 1–8.
- [110] X. Li and J. Li, “AoE: Angle-optimized embeddings for semantic textual similarity,” in *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Bangkok, Thailand: Association for Computational Linguistics, 2024, pp. 1825–1839, previously known as “AngIE”. [Online]. Available: <https://aclanthology.org/2024.acl-long.101>
- [111] O. Muliarevych, “Enhancing system security: Llm-driven defense against prompt injection vulnerabilities,” in *2024 IEEE 17th International Conference on Advanced Trends in Radioelectronics, Telecommunications and Computer Engineering (TCSET)*, 2024, pp. 420–423.
- [112] B. Cao, Y. Cao, L. Lin, and J. Chen, “Defending against alignment-breaking attacks via robustly aligned LLM,” in *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, L.-W. Ku, A. Martins, and V. Srikumar, Eds. Bangkok, Thailand: Association for Computational Linguistics, Aug. 2024, pp. 10 542–10 560. [Online]. Available: <https://aclanthology.org/2024.acl-long.568/>
- [113] K. Huang, X. Liu, Q. Guo, T. Sun, J. Sun, Y. Wang, Z. Zhou, Y. Wang, Y. Teng, X. Qiu, Y. Wang, and D. Lin, “Flames: Benchmarking value alignment of LLMs in Chinese,” in *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, K. Duh, H. Gomez, and S. Bethard, Eds. Mexico City, Mexico: Association for Computational Linguistics, Jun. 2024, pp. 4551–4591. [Online]. Available: <https://aclanthology.org/2024.naacl-long.256/>



JAQUELINE D. DUARTE received her bachelor's degree in Informatics from the Federal University of Mato Grosso (UFMT) in 2010, along with a MBA in Big Data Analytics (2019). She is currently pursuing a master's degree in Electrical Engineering at the University of Brasília (UnB). She has over five years of experience in Data Analytics, Cyber Threat Intelligence and Threat Hunting and Red teaming. In 2022, she was recognized as a Global Leader by FS-ISAC LATAM for her significant contributions to threat information sharing. Her research interests include machine learning for cybersecurity and supply chain, AI governance and security.



GUILHERME CANDIDO received his bachelor's degree in Civil Engineering from the State University of Maringá (UEM) in 2013. He is currently pursuing a master's degree in Electrical Engineering at University of Brasília (UnB). He has over five years of experience in Open Source Intelligence, Penetration Testing and Red Teaming. His research interests include human factors in cybersecurity, social engineering, phishing, and spoofing attacks.



JOSÉ RICARDO A. DE BRITO FILHO received his bachelor's degree in Computer Science from University of Brasília (UnB) in 2010, his MBA degree in Data Science and Big Data from Fundação Getúlio Vargas (FGV) in 2019 and his MBA degree in Cybersecurity from Ibmec in 2025. He is currently pursuing a master's degree in Electrical Engineering at UnB and a law degree at Universidade Paulista (UNIP). He has worked as a systems analyst at the Electoral Court in Brasília since 2010. His research interests include security in embedded systems.



JOÃO SOUZA NETO PhD in Electrical Engineering from UnB, Master's in Electronic Engineering from Philips International Institute in the Netherlands. He is an Associate Researcher in the Master's in Cybersecurity program at the University of Brasília. He is certified as a CGEIT Trainer, CRISC Trainer, CDSPE Trainer, COBIT 2019 Trainer, COBIT 2019 Design Implementation, COBIT 5 Trainer, COBIT Certified Assessor, PMP, RMP, RCDD, and is an IEEE Senior Member.

He is a founding member of the Brasília Chapter of ISACA.



ELENA J. DA COSTA received her bachelor's degree from Islamic Azad University in 2011, her M.Sc. degree in biomedical engineering from Ilmenau University of Technology in 2019, and her Ph.D. degree in mechatronic systems from the University of Brasília (UnB) in 2024.

From 2019 to 2020, she worked as a Software Test Engineer in the automotive industry at Formel D, contributing to a BMW-related project. In 2024 and 2025, she was a postdoctoral project leader at

the University of Kassel. Her research interests include array signal processing, machine learning, and cybersecurity.



JOÃO PAULO JAVIDI DA COSTA received the Diploma degree in electronic engineering in 2003 from the Military Institute of Engineering (IME), Rio de Janeiro, Brazil, his M.Sc. degree in telecommunications in 2006 from the University of Brasília (UnB), Brazil, and his Ph.D. degree in electrical engineering in 2010 at Ilmenau University of Technology (TU Ilmenau), Germany.

Since August 2020, Prof. da Costa has been a professor of applied electrical engineering at the

Hamm-Lippstadt University of Applied Sciences in Germany and became a research professor at HSHL in March 2022. He is also a professor member of the Promotionskolleg NRW (PK NRW) for Ph.D. supervision.

He is a Senior Member of IEEE and has published more than 195 scientific papers and patents. He has received seven Best Paper Awards in international conferences. His research interests include autonomous vehicles, 6G, GNSS, and adaptive and array signal processing.



LAERTE PEOTTA DE MELO is the Digital Security Manager at one of largest Brazilian bank. He holds an M.Sc. (2008) and a Ph.D. (2012) in Electrical Engineering from the University of Brasília (UnB), along with a specialization in network security from the Catholic University of Brasília (2004).

With over 22 years of experience, he leads Red Team, black Team, Cyber Intelligence, and Threat Hunting teams. He also serves as an advisor in the

Master's Program in Electrical Engineering at UnB. He is the creator of BBCode, a transaction authorization and authentication protocol using visual 2FA, and holds three patents. His research focuses on cybercrime and digital fraud.

...